



# Advanced Distributed Systems

Lecture 08-Autoscaling

Dr. Zahra Najafabadi Samani



# Agenda



- Motivation
- Autoscaling
- Different type of autoscaling
- Autoscaling in Kubernetes



# Motivation


# Why autoscaling?

Workloads in various applications often exhibit significant fluctuations over time.



# Why autoscaling?

Workloads in various applications often exhibit significant fluctuations over time.



# Why autoscaling?



Scale up  
Increase the resources

Stay responsiveness  
Prevent down time  
Increase the performance

# Why autoscaling?



Scale down  
decrease the resources

Save cost  
Reduce resource wastage

# Why autoscaling?

Workloads in various applications often exhibit significant fluctuations over time leading to several critical challenges:

- **Resource inefficiency:** During periods of low demand system provides higher capacity than necessary, resulting in resource inefficiency such as escalated costs and unnecessary resource wastage.
- **QoS violation:** The system's incapacity to dynamically adjust to varying workloads can lead to performance issues such as increased response times and the risk of service failure, particularly during peak usage periods.
- **Increased management complexity:** Manually managing resources to meet fluctuating demand can be complex and time-consuming and making it challenging to efficiently allocate resources and respond quickly.



# Autoscaling



# What is autoscaling?

- Automatically adjusts computing resources based on real-time demand.
- Relies on automated processes and orchestration to reduce the need for manual intervention.
- Quickly scales resources in response to fluctuations in demand for consistent performance.
- Ensures optimal performance by scaling resources up or down to match current workloads.
- Controls costs by provisioning and utilizing resources only when necessary.
- Allows administrators to define policies based on metrics like CPU utilization, or other metrics.
- Integrated with load balancing to evenly distribute traffic and optimize resource utilization.

# Advantage of autoscaling

**Efficient resource utilization:** Automatically adjusts resources based on demand, minimizing underutilization or overprovisioning.

**Cost savings:** Ensures cost efficiency by scaling resources up during peak demand and down during periods of low activity, avoiding unnecessary expenses.

**Improved performance:** Maintains optimal performance levels by dynamically adapting to changes in workload, preventing performance degradation during high-demand periods.

**Enhanced scalability:** Facilitates the seamless expansion or contraction of infrastructure, allowing systems to handle increased loads without manual intervention.

**Reduced downtime:** Enhances system reliability by distributing workloads and adapting to changes, reducing the risk of service outages and downtime.

# Advantage of autoscaling

- **Automatic response to demand fluctuations:** Provides a quick and automated response to changes in demand, ensuring that resources align with the current needs of the system.
- **Increased flexibility:** Offers flexibility in managing varying workloads without requiring constant manual adjustments, allowing for a more adaptive and responsive infrastructure.



# Different type of autoscaling

# Different types of autoscaling

Fundamentally there are three types of autoscaling:


- Reactive
- Proactive
- Hybrid

# Reactive autoscaling

Reactive autoscaling involves dynamically adjusting resources based on immediate changes in demand, responding to predefined triggers or conditions rather than proactively anticipating future requirements.

- **Key features:**

- **Monitoring:** Continuously monitoring key performance indicators (KPIs) such as CPU usage, memory, and network traffic in real-time.
- **Scaling policies:** defining rules for resource allocation adjustments and reactively response to specific conditions.
- **Common Metrics:** CPU utilization, network traffic, or response times are often used as triggers for reactive scaling.
- **Triggered by Events:** triggering by predefined events or conditions, such as increased traffic or high resource utilization.



# Advantages of Reactive autoscaling

- **Simplicity:** It is often simpler to implement, especially with rule-based systems.
- **Cost Control:** In some scenarios, reactive autoscaling may help control costs by only scaling resources when predefined triggers are met. This can prevent unnecessary scaling during periods of low demand.
- **Predictability:** With well-defined triggers and policies, reactive autoscaling can provide a predictable response to specific events, making it easier to understand and manage.
- **Continuous Improvement:** Regularly review and refine the reactive autoscaling setup based on performance analysis and feedback. Adjust trigger conditions, scaling policies, and other parameters to enhance efficiency and responsiveness.

# Reactive autoscaling limitation

- **Latency:** One of the primary drawbacks is the inherent latency between the occurrence of a triggering event and the initiation of scaling actions. This delay can impact performance during sudden spikes in demand.
- **Overprovisioning/underprovisioning:** Reactive autoscaling may lead to overprovisioning, where resources are added too late or in excess. This can result in unnecessary costs and underutilized resources during periods of decreased demand.
- **Limited Adaptability:** Rule-based reactive systems may lack adaptability to dynamic and unpredictable workloads. These systems may struggle to cope with scenarios not covered by predefined rules.
- **Complexity in Rule Tuning:** Fine-tuning the rules for reactive autoscaling can be challenging. Constant adjustments may be needed to avoid both false positives and negatives, making it an ongoing maintenance task.

# Proactive autoscaling

Proactive autoscaling employs predictive analytics and machine learning to forecast upcoming changes in demand. It dynamically adjusts resources ahead of actual spikes, optimizing performance and resource utilization.

- **Key features:**
  - **Monitoring and Analytics:**
    - Continuous monitoring of key performance indicators (KPIs).
  - **Predictive Algorithms:**
    - Utilization of historical data for trend analysis.
    - Machine learning models to forecast future demand.
    - Identify patterns and anomalies.
  - **Automated Scaling:**
    - Trigger scaling actions before reaching critical points.
    - Preemptive adjustments based on predictions.

# Advantage of proactive autoscaling

- **Improved Performance:** Anticipating future demand allows for timely resource allocation, ensuring optimal system performance and responsiveness during peak periods.
- **Cost Optimization:** Proactive autoscaling helps in efficient resource utilization by scaling up or down based on predicted demand. This can lead to cost savings by avoiding unnecessary overprovisioning.
- **Scalability:** Proactive autoscaling is well-suited for applications with rapidly changing workloads, enabling systems to scale seamlessly in response to evolving usage patterns.
- **Predictive Analytics:** Leveraging machine learning and predictive analytics allows for more accurate forecasting, enabling organizations to make informed decisions about resource allocation.
- **Reduced Operational Burden:** With the ability to automatically adjust resources based on predictions, proactive autoscaling reduces the need for manual intervention and constant monitoring, easing the operational burden on administrators.


# Proactive autoscaling limitation

- **Complex Implementation:** Implementing predictive analytics and machine learning models can be complex and may require expertise in data science.
- **Potential for Inaccuracies:** Predictive models may not always accurately forecast future demand, leading to overestimation or underestimation of resource needs.
- **Resource Costs for Predictive Analytics:** Developing and maintaining predictive analytics models may incur additional resource costs.
- **Adaptability to Rapid Changes:** Proactive autoscaling may face challenges in adapting to rapid and unforeseen changes in demand that fall outside the scope of historical patterns.
- **Continuous Model Refinement:** Predictive models require ongoing refinement and adjustment to accommodate changes in application behavior, making it a continuous process that demands attention and resources.

# Hybrid autoscaling

Hybrid autoscaling combines both reactive and proactive scaling strategies to create a dynamic and flexible approach to managing resources in response to changing demands.

Hybrid autoscaling approaches leverages the strengths of both reactive and proactive scaling mechanisms to achieve optimal performance, resource utilization, and cost efficiency.



# Different types of autoscaling

Auto-scaling encompasses a variety of strategies tailored to meet the dynamic demands of applications. Three primary types include:


- Horizontal autoscaling
- Vertical autoscaling
- Hybrid autoscaling



# Horizontal autoscaling

Horizontal auto-scaling involves adding or removing identical instances of resources, such as servers, to accommodate changing workloads. It is especially important for companies that need **high availability** services with a requirement for minimal downtime.

- Key Points:
  - **Scalability:** Achieved by adding more instances to the system.
  - **Load Balancing:** Traffic is distributed across multiple instances to ensure even resource usage.




# Advantages of horizontal autoscaling

- **Increased High Availability:**
  - Spreading infrastructure across multiple areas enhances availability.
  - If one machine fails, others can seamlessly handle the load.
- **Reduced Downtime:**
  - Adding machines allows scaling without the need to switch off the old machine.
  - Minimal or no downtime required for effective scaling.
- **Easy Resizing:**
  - Scalability allows effortless adjustment according to current needs.
- **Cost Efficiency:**
  - Cost can be linked to usage and you don't always have to pay for peak demand.

# Horizontal autoscaling limitation


- **Use case limitation:**
  - Applications must be designed to scale horizontally.
- **Increasing complexity:**
  - Shared state and session management can be challenging.
- **Increasing Initial costs:**
  - Adding new servers is far more expensive than upgrading old ones.



# Vertical autoscaling

Vertical auto-scaling Involves adjusting the capacity of a single resource, such as increasing the CPU, RAM, or storage of a server. The expansion has an upper limit based on the capacity of the server or machine being expanded.

- Key Points:
  - **Scaling Up:** Adding more power to existing instances by increasing resources.
  - **Scaling Down:** Reducing resources to optimize performance and cost.



# Advantages of vertical autoscaling

- **Cost-effective:**
  - Upgrading a pre-existing server costs less than purchasing a new one. It is less likely to add new backup and virtualization software when scaling vertically. Maintenance costs may potentially remain the same too.
- **Less complex process communication:**
  - When a single node handles all the layers of your services, it will not have to synchronize and communicate with other machines to work. This may result in faster responses.
- **Less complicated maintenance:**
  - Not only the maintenance is cheaper but it is less complex because of the number of nodes you will need to manage.
- **Less need for software changes:**
  - It is less likely to change how the software on a server works or how it is implemented.

# Vertical autoscaling limitation

- **Downtime Considerations:**
  - Without a backup server capable of handling operations, downtime is inevitable during machine upgrades.
- **Risk of Single Point of Failure:**
  - Potential loss of all data in the event of a hardware or software failure.
- **Upgrade Constraints:**
  - Limited by the maximum capacity of a single resource.



# Horizontal vs Vertical autoscaling

	Horizontal scaling	Vertical scaling
Description	Increase or decrease the number of nodes in a cluster or system to handle an increase or decrease in workload	Increase or decrease the power of a system to handle increased or reduced workload
Example	Add or reduce the number of virtual machines (VM) in a cluster of VMs	Add or reduce the CPU or memory capacity of the existing VM
Execution	Scale in/out	Scale up/down
Workload distribution	Workload is distributed across multiple nodes.	A single node handles the entire workload.
Required architecture	Distributed	Any

# Horizontal vs Vertical autoscaling

	Horizontal scaling	Vertical scaling
<b>Complexity and maintenance</b>	Higher	Lower
<b>Implementation</b>	Takes more time, expertise, and effort	Takes less time, expertise, and effort
<b>Configuration</b>	This requires modifying a sequential piece of logic in order to run workloads concurrently on multiple machines	No need to change the logic. The same code can run on a higher-spec device
<b>Costs</b>	High costs initially; optimal over time	Low-cost initially; less cost-effective over time
<b>Networking</b>	Slower machine-to-machine communication	inter-machine communication
<b>Performance</b>	Higher	Lower

# Horizontal vs Vertical autoscaling

	Horizontal scaling	Vertical scaling
Limitation	Add as many machines as you can	Limited to the resource capacity the single machine can handle



# Autoscaling in Kubernetes

# Autoscaling in Kubernetes


Kubernetes provides three types of autoscaling:

# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:

- Horizontal pod autoscaling (HPA):

Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.




# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:

- Horizontal pod autoscaling (HPA):

Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.




# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:

- Horizontal pod autoscaling (HPA):



Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.



# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:


- Horizontal pod autoscaling (HPA):  
Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.
- Vertical pod autoscaling (VPA)  
Automatically adjusts the resource requests (CPU and memory) of individual containers within a pod based on observed usage.



# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:


- Horizontal pod autoscaling (HPA):  
Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.
- Vertical pod autoscaling (VPA)  
Automatically adjusts the resource requests (CPU and memory) of individual containers within a pod based on observed usage.



# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:



- Horizontal pod autoscaling (HPA):  
Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.
- Vertical pod autoscaling (VPA)  
Automatically adjusts the resource requests (CPU and memory) of individual containers within a pod based on observed usage.



# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:




- Horizontal pod autoscaling (HPA):  
Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.
- Vertical pod autoscaling (VPA)  
Automatically adjusts the resource requests (CPU and memory) of individual containers within a pod based on observed usage.
- Cluster autoscaling (CA):  
Automatically adjust the number of nodes based on the resource requirements and constraints of the workloads running in the cluster.



# Autoscaling in Kubernetes

Kubernetes provides three types of autoscaling:




- Horizontal pod autoscaling (HPA):  
Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.
- Vertical pod autoscaling (VPA)  
Automatically adjusts the resource requests (CPU and memory) of individual containers within a pod based on observed usage.
- Cluster autoscaling (CA):  
Automatically adjust the number of nodes based on the resource requirements and constraints of the workloads running in the cluster.



# Autoscaling in Kubernetes


Kubernetes provides three types of autoscaling:

- Horizontal pod autoscaling (HPA):  
Automatically adjusts the number of pod replicas in a deployment or replica set based on observed metrics, such as CPU utilization or custom metrics.
- Vertical pod autoscaling (VPA)  
Automatically adjusts the resource requests (CPU and memory) of individual containers within a pod based on observed usage.
- Cluster autoscaling (CA):  
Automatically adjust the number of nodes based on the resource requirements and constraints of the workloads running in the cluster.



# Horizontal Pod Autoscaling

- Kubernetes monitors the specified metrics for each pod in the target deployment.
- If the observed metrics go beyond or fall below the defined thresholds, the HPA controller triggers scaling actions.
- Scaling actions involve adjusting the number of pod replicas to meet the desired metric targets.




# Horizontal Pod Autoscaling





```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: utility-api
spec:
  minReplicas: 1
  maxReplicas: 5
  metrics:
    - resource:
        name: cpu
        target:
          averageUtilization: 70
          type: Utilization
        type: Resource
    scaleTargetRef:
      apiVersion: apps/v1
      kind: Deployment
      name: utility-api
```

Master Node





Worker Node


# Horizontal Pod Autoscaling




# Horizontal Pod Autoscaling




# Horizontal Pod Autoscaling




# Horizontal Pod Autoscaling




# Horizontal Pod Autoscaling



# Horizontal Pod Autoscaling



# Horizontal Pod Autoscaling



# Horizontal Pod Autoscaling

Calculating the desired number of replicas:

$$\text{desiredReplicas} = \left\lceil \text{currentReplicas} \cdot \frac{\text{currentMetricValue}}{\text{desiredMetricValue}} \right\rceil$$

# Horizontal Pod Autoscaling

Calculating the desired number of replicas:

$$\text{desiredReplicas} = \left\lceil \text{currentReplicas} \cdot \frac{\text{currentMetricValue}}{\text{desiredMetricValue}} \right\rceil$$

Example:

**CPU:**  $\text{desiredReplicas} = \left\lceil 2 \cdot \frac{90}{70} \right\rceil = 3 \text{ Replica}$

**Memory:**  $\text{desiredReplicas} = \left\lceil 2 \cdot \frac{1000}{500} \right\rceil = 4 \text{ Replica}$

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: utility-api
spec:
  minReplicas: 1
  maxReplicas: 5
  metrics:
    - resource:
        name: cpu
        target:
          averageUtilization: 70
        type: Utilization
      type: Resource
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: utility-api
```

# Horizontal Pod Autoscaling


Calculating the desired number of replicas:

$$\text{desiredReplicas} = \left\lceil \text{currentReplicas} \cdot \frac{\text{currentMetricValue}}{\text{desiredMetricValue}} \right\rceil$$

Example:

**CPU:** desiredReplicas =  $\left\lceil 2 \cdot \frac{90}{70} \right\rceil = 3 \text{ Replica}$

**Memory:** desiredReplicas =  $\left\lceil 2 \cdot \frac{1000}{500} \right\rceil = 4 \text{ Replica}$



```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: utility-api
spec:
  minReplicas: 1
  maxReplicas: 5
  metrics:
    - resource:
        name: cpu
        target:
          averageUtilization: 70
        type: Utilization
      type: Resource
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: utility-api
```

# Horizontal Pod Autoscaling

Creating a deployment:

```
! deployment.yaml ×  
autoscaling > ! deployment.yaml > {} spec > {} template > {} spec  
  io.k8s.api.apps.v1.Deployment (v1@deployment.json)  
  1 apiVersion: apps/v1  
  2 kind: Deployment  
  3 metadata:  
  4   name: utility-api  
  5 spec:  
  6   replicas: 2  
  7   selector:  
  8     matchLabels:  
  9       app: utility-api  
10   template:  
11     metadata:  
12       name: utility-api-pod  
13       labels:  
14         app: utility-api  
15     spec:  
16       containers:  
17         - name: utility-api  
18           image: pavanelthepu/utility-api  
19           ports:  
20             - containerPort: 8080  
21           resources:  
22             requests:  
23               memory: 20Mi
```

Creating a service to access the deployment:

```
! deployment.yaml      ! service.yaml ×  
autoscaling > ! service.yaml > {} spec > [ ]ports  
  io.k8s.api.core.v1.Service (v1@service.json)  
  1 apiVersion: v1  
  2 kind: Service  
  3 metadata:  
  4   name: utility-api-service  
  5 spec:  
  6   selector:  
  7     app: utility-api  
  8   ports:  
  9     - port: 8080  
10       targetPort: 8080
```

# Horizontal Pod Autoscaling

Generating HPA for the deployment:

```
! deployment.yaml      ! hpa.yaml  X
autoscaling > ! hpa.yaml > {} spec > [ ]metrics > {} 0 > {}
1   apiVersion: autoscaling/v2
2   kind: HorizontalPodAutoscaler
3   metadata:
4     name: utility-api
5   spec:
6     minReplicas: 1
7     maxReplicas: 5
8     metrics:
9       - resource:
10         name: cpu
11         target:
12           averageUtilization: 70
13           type: Utilization
14         type: Resource
15     scaleTargetRef:
16       apiVersion: apps/v1
17       kind: Deployment
18       name: utility-api
```

# Horizontal Pod Autoscaling

Automatically scaling down the number of pod replicas:

```
avan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
utility-api-6cd64bbb76-4vbxr  1/1     Running   0          3m16s
utility-api-6cd64bbb76-8vmj7  1/1     Running   0          3m16s
avan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
utility-api-6cd64bbb76-4vbxr  2m           89Mi
avan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
utility-api-6cd64bbb76-4vbxr  1/1     Running   0          7m28s
avan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get hpa
NAME          REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
utility-api   Deployment/utility-api  1%/70%       1          5          1          9m25s
avan.elthepu@rrcs-172-254-255-63 autoscaling %
```

Number of replicas  
before autoscaling

# Horizontal Pod Autoscaling

Automatically scaling down the number of pod replicas:

```
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
utility-api-6cd64bbb76-4vbxr  1/1     Running   0          3m16s
utility-api-6cd64bbb76-8vmj7  1/1     Running   0          3m16s
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
utility-api-6cd64bbb76-4vbxr  2m          89Mi
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
utility-api-6cd64bbb76-4vbxr  1/1     Running   0          7m28s
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get hpa
NAME          REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
utility-api   Deployment/utility-api  1%/70%       1          5          1          9m25s
savan.elthepu@rrcs-172-254-255-63 autoscaling %
```

Number of replicas before autoscaling

Pod (service) utilization

# Horizontal Pod Autoscaling

Automatically scaling down the number of pod replicas:

```
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
utility-api-6cd64bbb76-4vbxr  1/1     Running   0          3m16s
utility-api-6cd64bbb76-8vmj7  1/1     Running   0          3m16s
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
utility-api-6cd64bbb76-4vbxr  2m          89Mi
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
utility-api-6cd64bbb76-4vbxr  1/1     Running   0          7m28s
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl get npa
NAME          REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
utility-api  Deployment/utility-api  1%/70%       1          5          1          9m25s
savan.elthepu@rrcs-172-254-255-63 autoscaling %
```

Number of replicas before autoscaling

Pod (service) utilization

Number of replicas after autoscaling


# Horizontal Pod Autoscaling

```
avan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl describe hpa utility-api
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPodAutoscaler
  name: utility-api
  namespace: default
  labels: <none>
  annotations: <none>
  creationTimestamp: Sat, 18 Feb 2023 19:47:12 +0530
  reference: Deployment/utility-api
    ( current / target )
  metrics:
    resource cpu on pods (as a percentage of request): 0% (2m) / 70%
  in replicas: 1
  max replicas: 5
  deployment pods: 1 current / 1 desired
  conditions:
    Type Status Reason Message
    ---- ----
    AbleToScale True ScaleDownStabilized recent recommendations were higher than current one, applying the highest recent recommendation
    ScalingActive True ValidMetricFound the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
    ScalingLimited False DesiredWithinRange the desired count is within the acceptable range
  events:
    Type Reason Age From Message
    ---- ----
    Warning FailedComputeMetricsReplicas 7m12s (x12 over 9m57s) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), error is: failed to get cpu resource metric value: failed to get cpu utilization: unable to get metrics for resource cpu: unable to get metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)
    Warning FailedGetResourceMetric 6m57s (x13 over 9m57s) horizontal-pod-autoscaler failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource metrics API: the server could not find the requested resource (get metrics.k8s.io)
    Normal SuccessfulRescale 4m57s horizontal-pod-autoscaler New size: 1; reason: All metrics below target
avan.elthepu@rrcs-172-254-255-63 autoscaling %
```

# Horizontal Pod Autoscaling

Increasing the resource utilization of pod(service):

1. Generating workload stress:



The screenshot shows a portion of a Java code editor with the following code:

```
@RestController
@RequestMapping("/api")
@CrossOrigin
public class UtilityController {

    private static final Logger logger = LoggerFactory.getLogger(UtilityController.class);

    @GetMapping("/stress")
    public int generateStress() {
        logger.info("Generating stress..");
        int counter = 0;
        for(int i = 0; i <= 5000000; i++) {
            counter += i;
        }
        return counter;
    }
}
```

A red arrow points from the line `for(int i = 0; i <= 5000000; i++) {` to a callout box containing the text "Create load over the CPU".

Create load over the CPU

# Horizontal Pod Autoscaling

Increasing the resource utilization of pod(service):

2. Creating alpine pod to communicate with utility api:

```
! deployment.yaml      ! traffic-generator.yaml X
autoscaling > ! traffic-generator.yaml > {} spec > [ ]c
          io.k8s.api.core.v1.Pod (v1@pod.json)
1   apiVersion: v1
2   kind: Pod
3   metadata:
4     name: traffic-generator
5   spec:
6     containers:
7       - name: alpine
8         image: alpine
9         args:
10        - sleep
11        - "100000000"
```

# Horizontal Pod Autoscaling

Increasing the resource utilization of pod(service):

3. Generate load on utility/ api pod so that the CPU utilization goes beyond 70% and we can see how pod scale up automatically. For that call stress api continuously so that the CPU usage goes up with using a http benchmarking tool called wrk\*.

```
savan.elthepu@rrcs-172-254-255-63 autoscaling % kubectl exec -it traffic-generator -- sh
$ # apk add wrk
fetch https://dl-cdn.alpinelinux.org/alpine/v3.17/main/x86_64/APKINDEX.tar.gz
fetch https://dl-cdn.alpinelinux.org/alpine/v3.17/community/x86_64/APKINDEX.tar.gz
1/3) Installing libgcc (12.2.1_git20220924-r4)
2/3) Installing luajit (2.1_p20210510-r3)
3/3) Installing wrk (4.1.0-r5)
Executing busybox-1.35.0-r29.trigger
OK: 9 MiB in 18 packages
$ # wrk -c 5 -t 5 -d 300s -H "Connection: Close" http://utility-api-service:8080/api/stress
```

Enter to the pod

\*<https://github.com/wg/wrk>

# Horizontal Pod Autoscaling

Increasing the resource utilization of pod(service):

3. Generate load on utility/ api pod so that the CPU utilization goes beyond 70% and we can see how pod scale up automatically. For that call stress api continuously so that the CPU usage goes up with using a http benchmarking tool called wrk\*.

Install wrk

```
bayan.elthebu@rrcs-172-254-255-63 autoscaling % kubectl exec -it traffic-generator -- sh  
1 # apk add wrk  
fetch https://dl-cdn.alpinelinux.org/alpine/v3.17/main/x86_64/APKINDEX.tar.gz  
fetch https://dl-cdn.alpinelinux.org/alpine/v3.17/community/x86_64/APKINDEX.tar.gz  
1/3) Installing libgcc (12.2.1_git20220924-r4)  
2/3) Installing luajit (2.1_p20210510-r3)  
3/3) Installing wrk (4.1.0-r5)  
Executing busybox-1.35.0-r29.trigger  
OK: 9 MiB in 18 packages  
# wrk -c 5 -t 5 -d 300s -H "Connection: Close" http://utility-api-service:8080/api/stress
```

Enter to the pod

\*<https://github.com/wg/wrk>

# Horizontal Pod Autoscaling

Increasing the resource utilization of pod(service):

3. Generate load on utility/ api pod so that the CPU utilization goes beyond 70% and we can see how pod scale up automatically. For that call stress api continuously so that the CPU usage goes up with using a http benchmarking tool called wrk\*.

Install wrk

```
# apk add wrk
fetch https://dl-cdn.alpinelinux.org/alpine/v3.17/main/x86_64/APKINDEX.tar.gz
fetch https://dl-cdn.alpinelinux.org/alpine/v3.17/community/x86_64/APKINDEX.tar.gz
1/3) Installing libgcc (12.2.1_git20220924-r4)
2/3) Installing luajit (2.1_p20210510-r3)
3/3) Installing wrk (4.1.0-r5)
Executing busybox-1.35.0-r29.trigger
OK: 9 MiB in 18 packages
# wrk -c 5 -t 5 -d 30s -H "Connection: Close" http://utility-api-service:8080/api/stress
```

Enter to the pod

Number of connections

Number of threads

Duration

Simulating the load

\*<https://github.com/wg/wrk>

# Horizontal Pod Autoscaling

Automatically scaling up the number of pod replicas:


```
pavan.elthepu@Pavan-Elthepu autoscaling % kubectl top pods
NAME                           CPU(cores)   MEMORY(bytes)
traffic-generator               1m           6Mi
utility-api-6cd64bbb76-4vbxr   43m          102Mi
pavan.elthepu@Pavan-Elthepu autoscaling % kubectl top pods
NAME                           CPU(cores)   MEMORY(bytes)
traffic-generator               57m          7Mi
utility-api-6cd64bbb76-4vbxr   301m         122Mi
utility-api-6cd64bbb76-snxfj   388m         106Mi
utility-api-6cd64bbb76-tk8v8   342m         104Mi
utility-api-6cd64bbb76-z9ct6   339m         104Mi
utility-api-6cd64bbb76-zmw6g   469m         108Mi
pavan.elthepu@Pavan-Elthepu autoscaling %
```

# Horizontal Pod Autoscaling


```
pavan.elthepu@Pavan-Elthepu autoscaling % kubectl describe hpa utility-api
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPo
scaler
  name: utility-api
  namespace: default
  labels: <none>
  annotations: <none>
  creationTimestamp: Sat, 18 Feb 2023 19:47:12 +0530
  reference: Deployment/utility-api
  metrics:
    resource cpu on pods (as a percentage of request): 117% (294m) / 70%
  min replicas: 1
  max replicas: 5
  deployment pods: 5 current / 5 desired
  conditions:
    Type Status Reason Message
    ---- ----
    AbleToScale True ReadyForNewScale recommended size matches current size
    ScalingActive True ValidMetricFound the HPA was able to successfully calculate a replica count from cpu resource utilization
    (percentage of request)
    ScalingLimited True TooManyReplicas the desired replica count is more than the maximum replica count
  events:
    Type Reason Age From Message
    ---- ----
    Warning FailedComputeMetricsReplicas 20m (x12 over 23m) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), first
    is: failed to get cpu resource metric value: failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch
    metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)
    Warning FailedGetResourceMetric 20m (x13 over 23m) horizontal-pod-autoscaler failed to get cpu utilization: unable to get
    metrics for resource cpu: unable to fetch metrics from resource metrics API: the server could not find the requested resource (get pods.
    .k8s.io)
    Normal SuccessfulRescale 18m horizontal-pod-autoscaler New size: 1; reason: All metrics below target
    Normal SuccessfulRescale 3m23s horizontal-pod-autoscaler New size: 4; reason: cpu resource utilization
    (percentage of request) above target
    Normal SuccessfulRescale 3m8s horizontal-pod-autoscaler New size: 5; reason: cpu resource utilization
```

# Vertical Pod Autoscaling

- The VPA consists of three components:
  - **Admission controller:** Triggers when new pods are created and checks if it is covered by a VPA.
  - **Recommender:** Provides advice based on historical and live resource usage metrics.
  - **Updater:** Decides which pods should be restarted based on resources allocation recommendation calculated by the Recommender.



# Vertical Pod Autoscaling



# Horizontal Pod Autoscaling

```
autoscaling > ! vpa.yaml > {} spec > {} targetRef :  
1   apiVersion: autoscaling.k8s.io/v1  
2   kind: VerticalPodAutoscaler  
3   metadata:  
4     name: utility-api  
5   spec:  
6     targetRef:  
7       apiVersion: apps/v1  
8       kind: Deployment  
9       name: utility-api  
10      updatePolicy:  
11        updateMode: "Off"
```

Auto  
Off  
Initial

Auto: It applies the recommendations that suggested by vertical autoscaler directly by updating the pods.

Off: It just gives you the recommendations but it does not update the replicas

Initial: It applies the recommended values only to the new created pod

Update mode: “off” is more recommended in productions because when vertical pod autoscaler applies the changes the pod will be restarted which might cause the workload disruption.

# Vertical Pod Autoscaling

```
Spec:  
  Target Ref:  
    API Version:  apps/v1  
    Kind:        Deployment  
    Name:        utility-api  
  Update Policy:  
    Update Mode: Off  
Status:  
  Conditions:  
    Last Transition Time: 2023-02-18T14:59:03Z  
    Status:          True  
    Type:           RecommendationProvided  
Recommendation:  
  Container Recommendations:  
    Container Name: utility-api  
    Lower Bound:  
      Cpu: 25m  
      Memory: 262144k  
    Target:  
      Cpu: 25m  
      Memory: 262144k  
    Uncapped Target:  
      Cpu: 25m  
      Memory: 262144k  
  Upper Bound:  
    Cpu: 100G  
    Memory: 100T
```



# Vertical Pod Autoscaling

- Install the vertical pod autoscaling

```
git clone http://github.com/Kubernetes/autoscaler.git  
cd autoscaler  
../vertical-pod-autoscaler/hack/vpa-up.sh
```

- Update the vertical pod autoscaler


```
Kubectl apply -f vpa.yaml
```

- Verify vertical pod autoscaler


```
pavan.elthepu@Pavan-Elthepu autoscaling % kubectl get vpa utility-api  
NAME      MODE   CPU    MEM     PROVIDED   AGE  
utility-api  Off    25m   262144k  True       74s  
pavan.elthepu@Pavan-Elthepu autoscaling % kubectl describe vpa utility-api
```

# Cluster Autoscaling

- The Cluster Autoscaler automatically adds or removes nodes in a cluster based on resource requests from pods.
- The Cluster Autoscaler doesn't directly measure CPU and memory usage values to make a scaling decision
- The Cluster Autoscaler checks every 10 seconds to detect any pods in a pending state, suggesting that the scheduler could not assign them to a node due to insufficient cluster capacity.



# Cluster Autoscaling



# Cluster Autoscaling limitation

- CA does not base scaling decisions on CPU or memory usage but solely on a pod's specified requests and limits for these resources.
- Unused computing resources requested by users are not identified by CA, leading to potential cluster waste and suboptimal utilization efficiency.
- When there's a need to scale up the cluster, CA initiates a scale-up request to the cloud provider within 30–60 seconds.
- The actual time taken by the cloud provider to create a new node can extend to several minutes or more.

Thank you for your attention 😊



Dr. Zahra Najafabadi Samani

Email: Zahra.Najafabadi-Samani@uibk.ac.at