universität innsbruck

Advanced Distributed Systems

Lecture 04-Prometheus and Grafana

Dr. Zahra Najafabadi Samani

Agenda

- Monitoring
- Prometheus
- Grafana
- Prometheus and Grafana on Kubernetes cluster

Monitoring

Monitoring

- Monitoring is one of the fundamental tasks in any system.
- It helps to detect problems and take action, or simply to know the current state of our systems.
- It helps to reduce (or even eliminate) downtime.
- It detects problems, and solves them before they impact users.

Monitoring

Monitoring helps us to:

- Expose Important **metrics (data)** over time
- Provide a way to configure **rules**, e.g (Cpu usage <80%)
- Send **alerts** (email) when rules are broken
- Provide **visualization/graphing**

Monitoring tools: Sensu, Prometheus, Nagios

Prometheus

Prometheus

- An open source monitoring and alerting system.
- Written in Go, fully published in 2015.
- An in-memory dimensional time series database.
- A custom query language.
- Designed to scrape metrics from application instances periodically based on service discovery.
- A data scraper that pulls metrics data over HTTP periodically at a configured interval.
- A time-series database to store all the metrics data.
- A simple user interface where you can visualize, query, and monitor all the metrics.

Prometheus Architecture

Prometheus server

- This is the heart of the system.
- It collects the metrics at regular intervals from multiple nodes and stores them locally.
- Prometheus primarily utilizes an HTTP endpoint to scrape metric data from applications. The metrics are either scraped by
  - the Prometheus server directly or
  - a push gateway for short-lived, instrumented jobs.

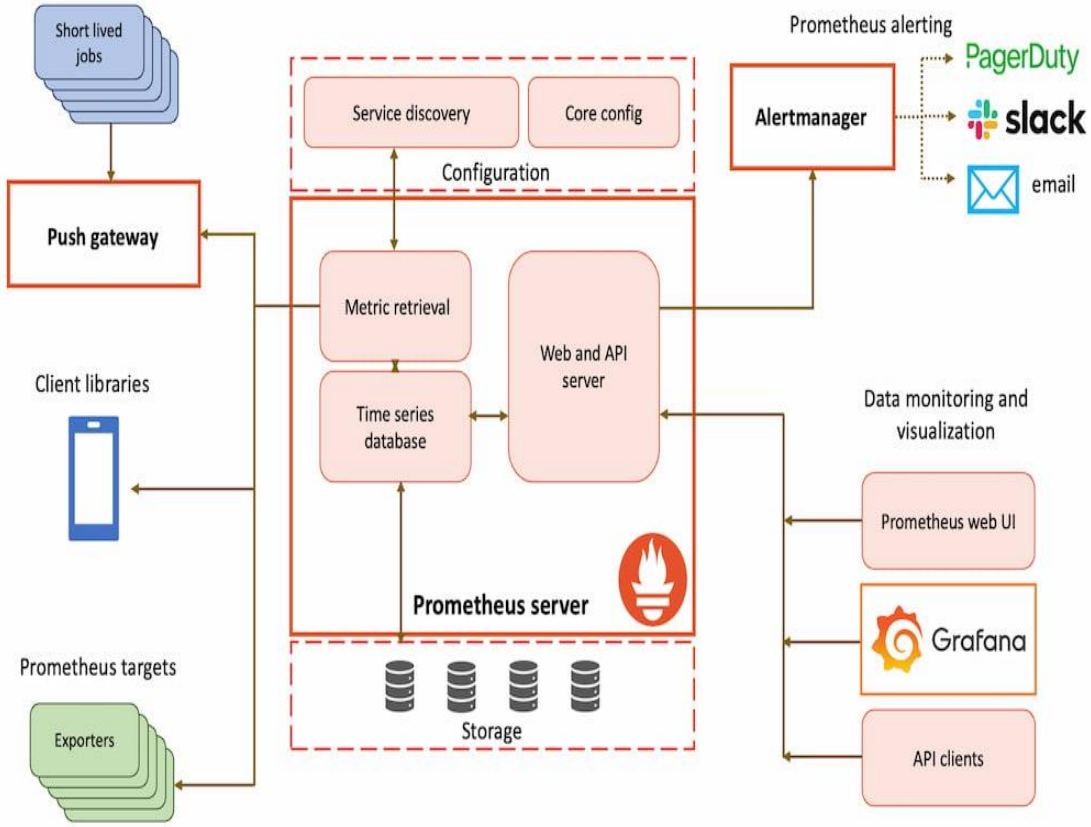

Prometheus Architecture

Client libraries

- Client libraries enable developers to add Prometheus instrumentation to their code.
- This instrumentation helps define and expose application metrics for each application language.
- Prometheus offers official libraries that are compatible with Ruby, Go, Python, Java/JVM, Haskell, Common Lisp, Dart, Erlang, C#/.NET, PHP, Perl, and Rust.

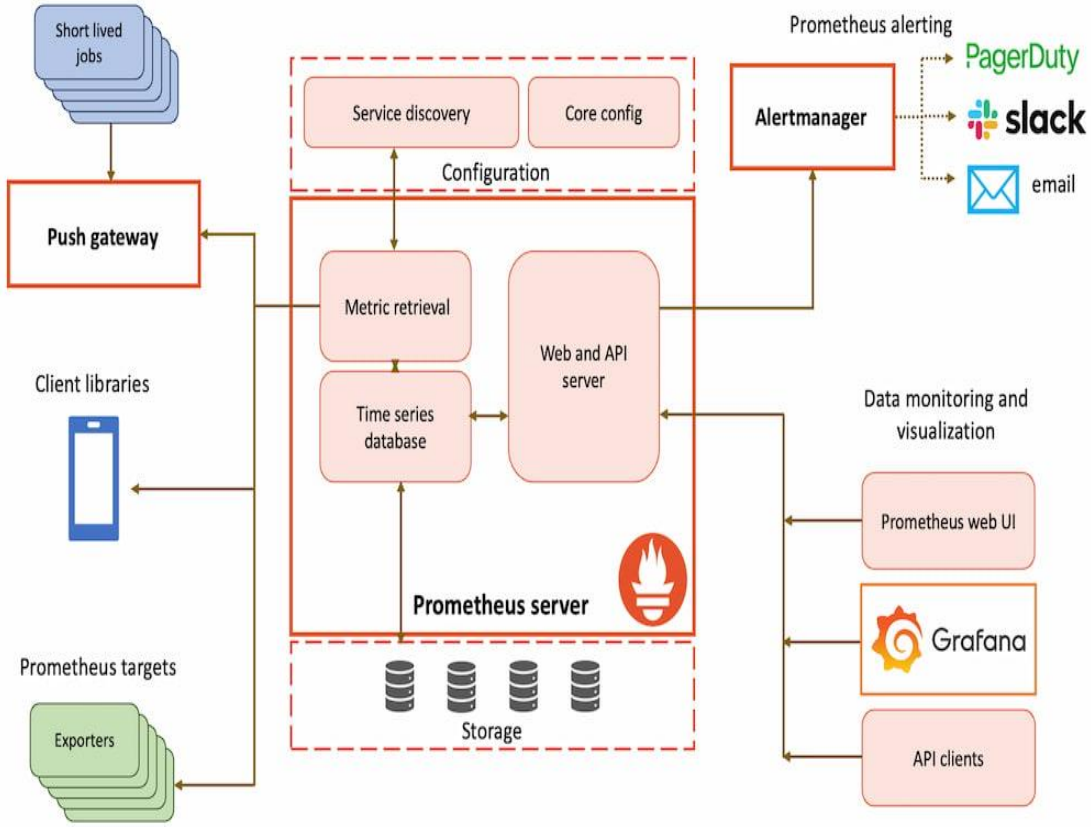

Prometheus Architecture

Exporter

- It refers to a separate software component or module that collects specific metrics from a system, service, or application.
- It acts as bridges between Prometheus and the target, translating the data from those systems into a format that Prometheus can understand and collect.

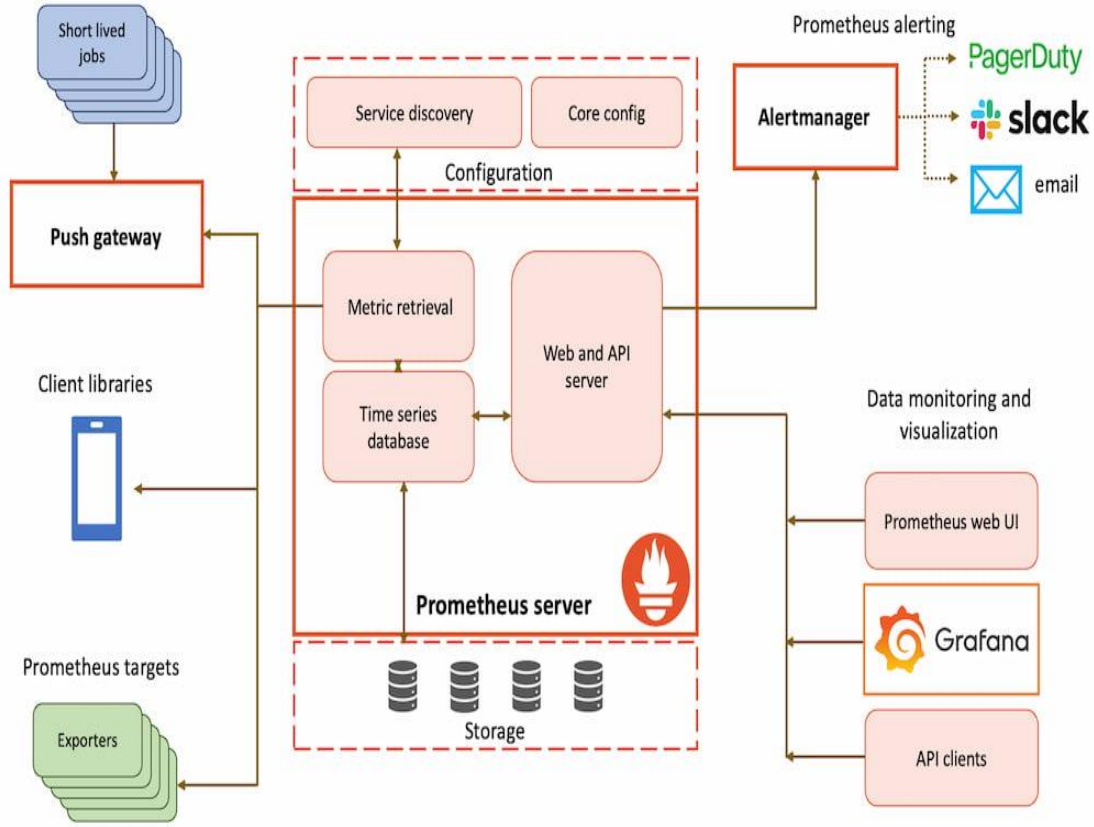

Prometheus Architecture

Service discovery

- Target discovery via static configuration or service discovery

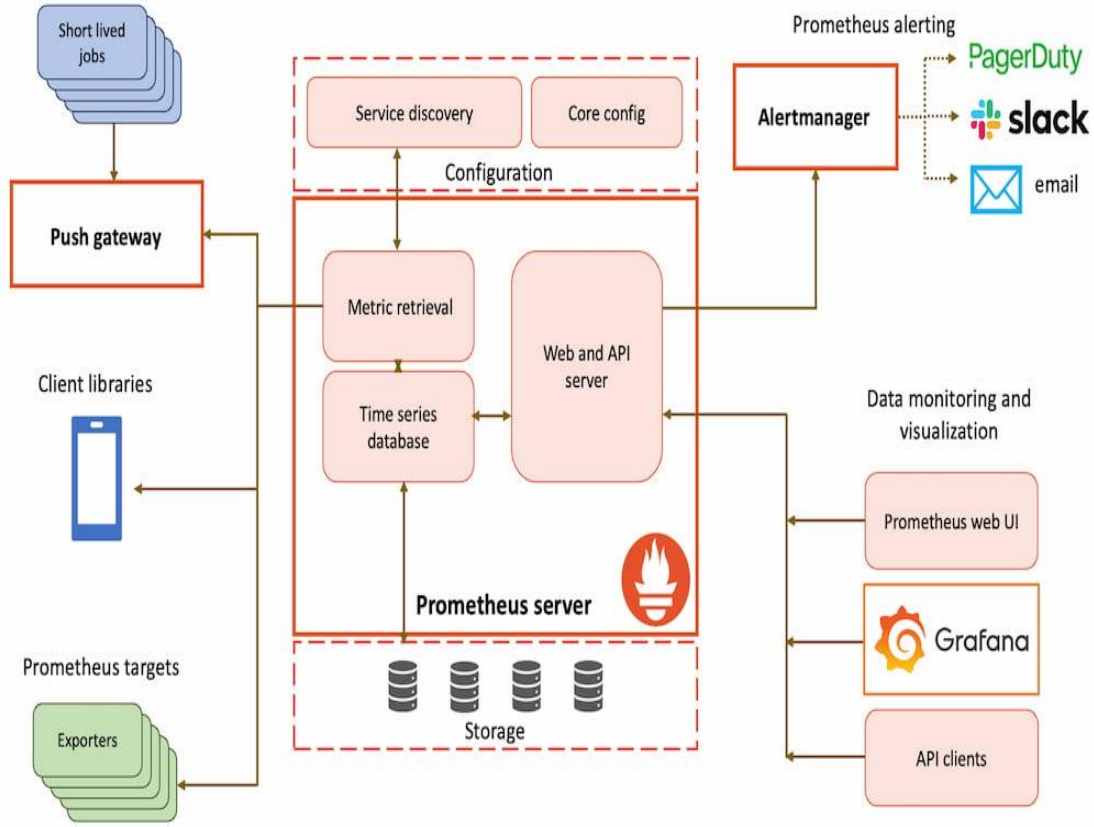

Prometheus Architecture

Visualization

- Prometheus comes with its own user interface to check on the configuration, nodes and graphs.
- It is compatible with Grafana, a leading open source visualization application, so that Prometheus data is available for viewing inside Grafana.
- It also exposes an API, so in case you are interested in writing your own clients.

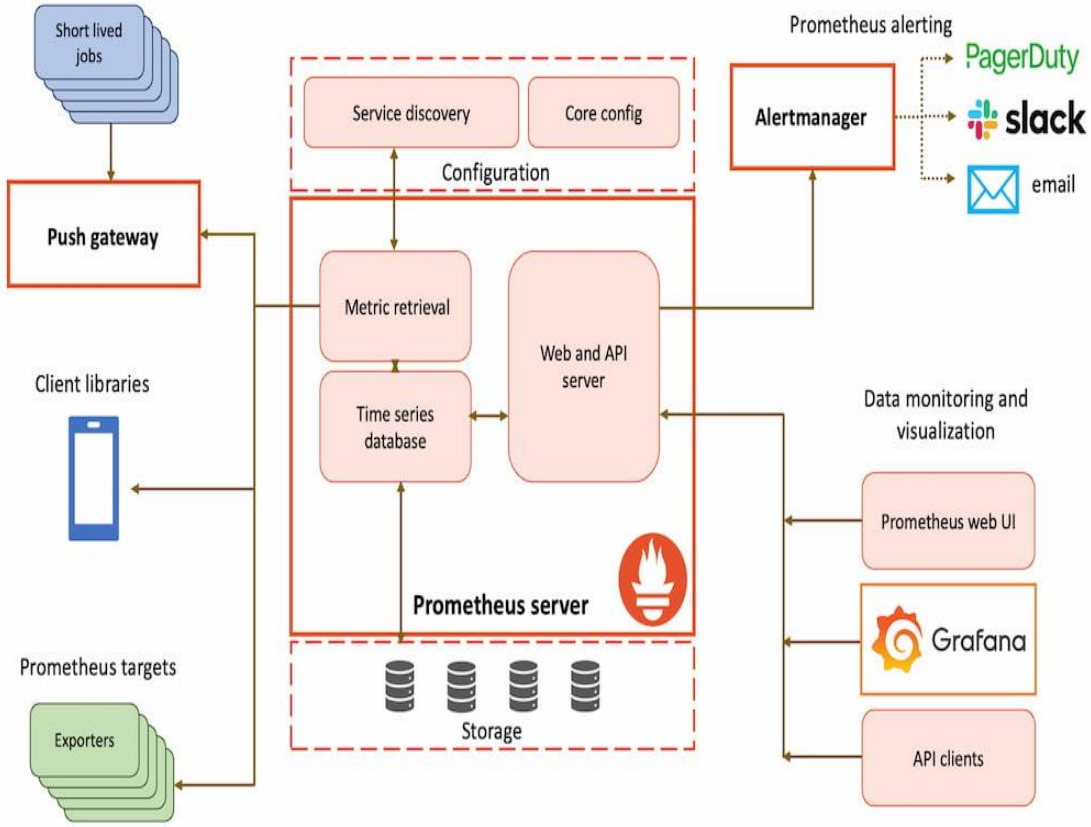

Prometheus Architecture

Alert manager

- AlertManager is a single binary which handles alerts sent by Prometheus server and notifies end user through E-mail, Slack or other tools.

Example (custom rules):

- The server is not running
- Average response time greater than 500ms
- The server's memory usage is greater than 95%.
- Number of 404 errors are greater than 10% of all requests.

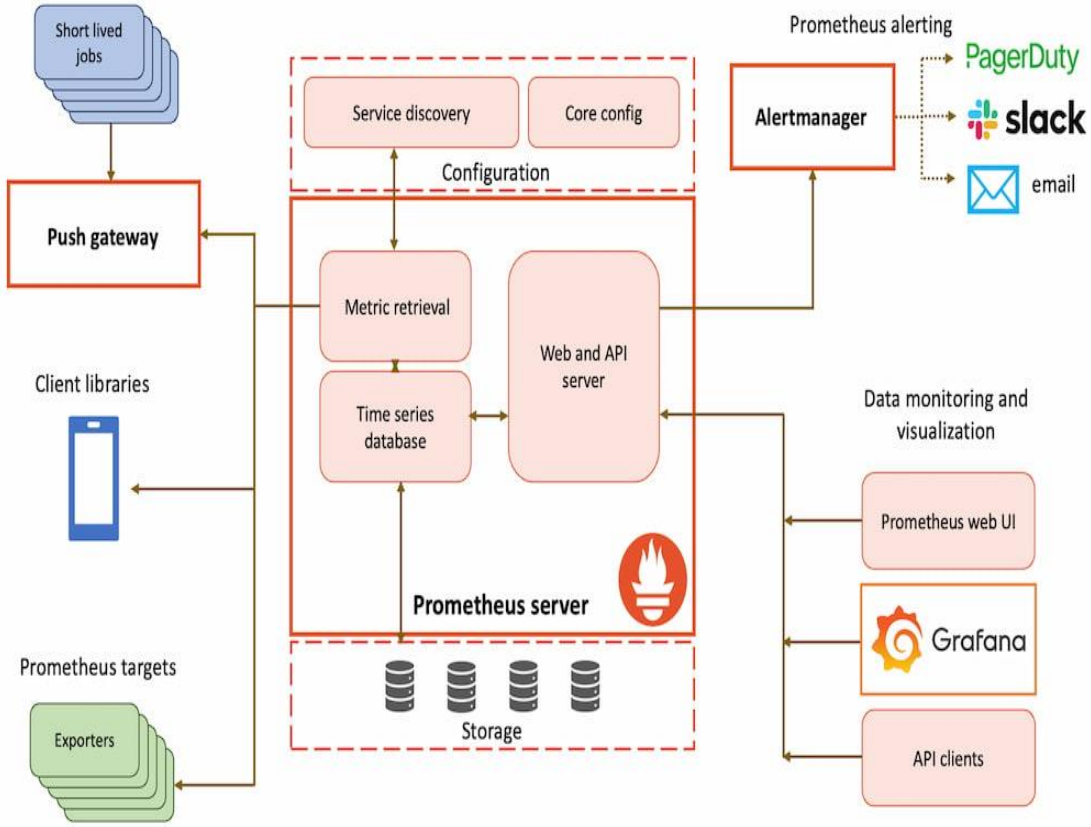

What Prometheus monitors

- Linux / windows servers (CPU status, memory, Disk space …)
- Apache server (On/Off status…)
- Single application (Exceptions, Requests count, client latency…)
- Service like database (status, Query Performance…)

Metrics

Metrics: all data as time series that Prometheus can scrape it and stores it in his database.

Metrics has two metadata lines: Help and Type

- **Help**: description of what the metrics is.
- **Type**: we have 3 metrics types:
  - Count: how many times x happened?
  - Gauge: what is the current value of x now?
  - Histogram; How long or how big?

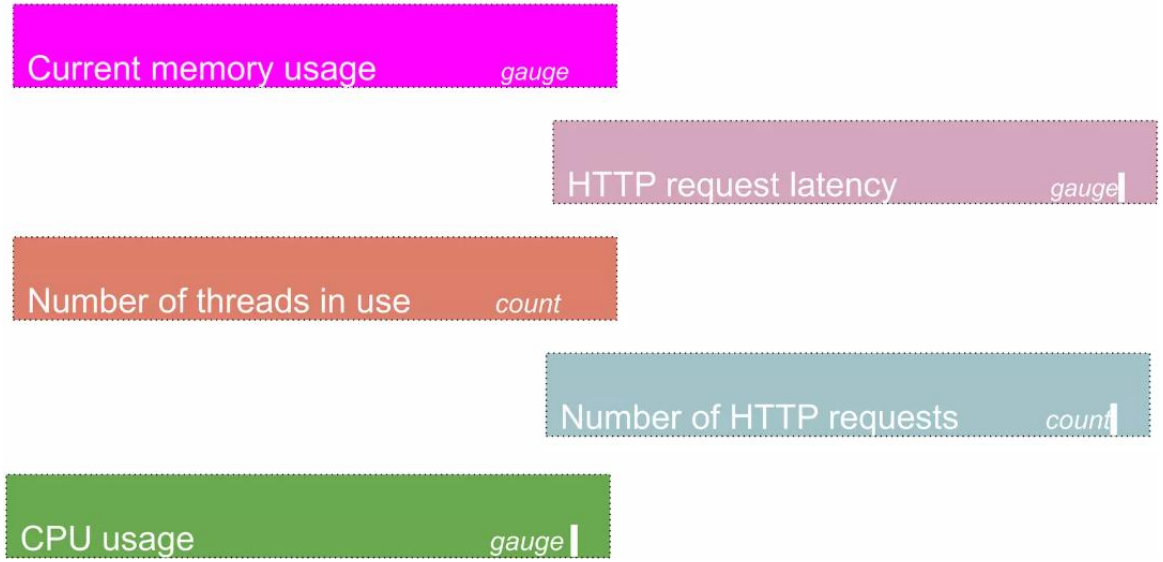

Prometheus query

- Prometheus provides a functional query language called PromQL (Prometheus Query Language) that lets the user select and aggregate time series data in real time.
- The result of an expression can either be shown as a graph, viewed as tabular data in Prometheus's expression browser, or consumed by external systems via the HTTP API.

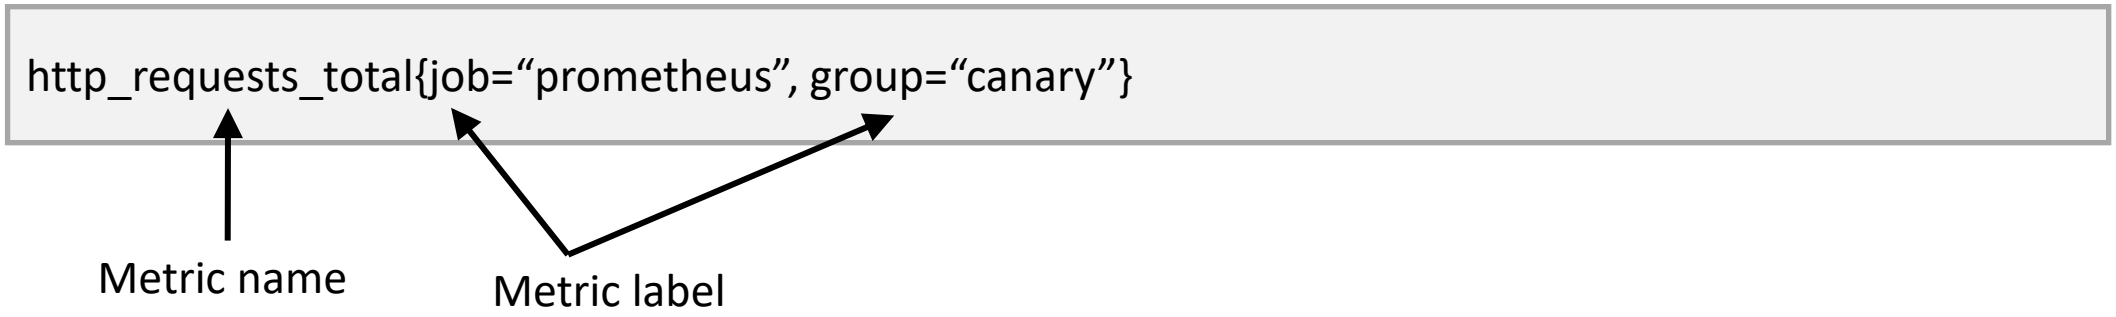

Prometheus query

Example Queries:

- http_requests_total:

  Return all time series with the metric

- http_requests_total{job="apiserver", handler="/api/comments"}:

  Return all time series with the metric http_requests_total and the given job and handler labels:

- Http_requests_total{status!~”4..”}:

  Query all HTTP status codes expect 4xx ones.

- Rate(Http_requests_total[5m])[30m:]:

  Returns the 5-minutes rate of the Http_requests_total metric for the past 30 min.

You can find more example here:
https://prometheus.io/docs/prometheus/latest/querying/examples/

Target endpoints and Exporter

In order to monitor anything with Prometheus, you need an exporter – a process that exposes metrics for Prometheus to harvest.

1. Prometheus initiates requests with parameters.
2. Probe to targets. The exporter subsequently starts the scrape after getting Prometheus' GET requests.
3. Return status code.
4. Return Prometheus metrics. Once the exporter is done with scraping, it will return the result in Prometheus metric format.

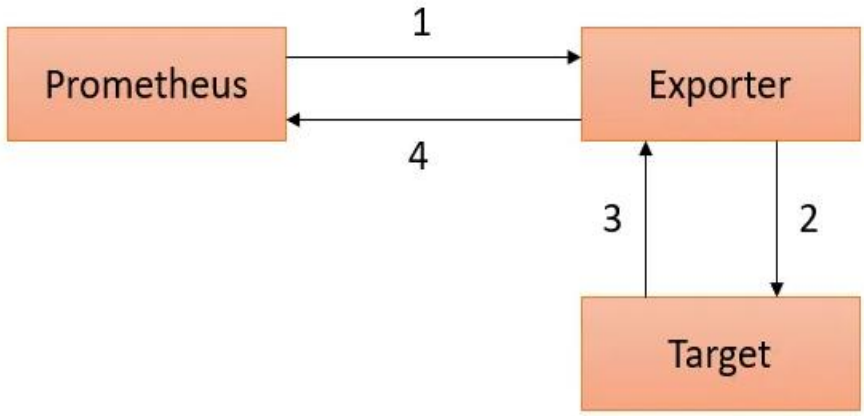

Target endpoints and Exporter

- **Global:**
  How often Prometheus will scrape its targets
- **Rule_files:**
  Rules for aggregating metric values or creating alerts when condition met
- **Scrape_configs:**
  What resources Prometheus monitors

```yaml
global:
  scrape_interval:     15s
  evaluation_interval: 15s

rule_files:
  # - "first.rules"
  # - "second.rules"

scrape_configs:
  - job_name: prometheus
    static_configs:
      - targets: ['localhost:9090']
```

prometheus.yml

Grafana

Grafana

- Grafana Is the open source analytics & monitoring solution for every database.
- Grafana allows you to bring data from various data sources like Elasticsearch, Prometheus, Graphite, InfluxDB etc, and visualize them with beautiful graphs.

Grafana features

- Create graphs easily
- Visualize data in different ways
- Show graphs for different time periods
- Create dashboard to be viewed by others
- Allows us to easily make changes

Prometheus and Grafana use cases

spring boot

NGINX

kubernetes

PostgreSQL

Microsoft SQL Server

We monitor all this targets using Prometheus and Grafana

Prometheus and Grafana on Kubernetes cluster

Prometheus and Grafana for Kubernetes

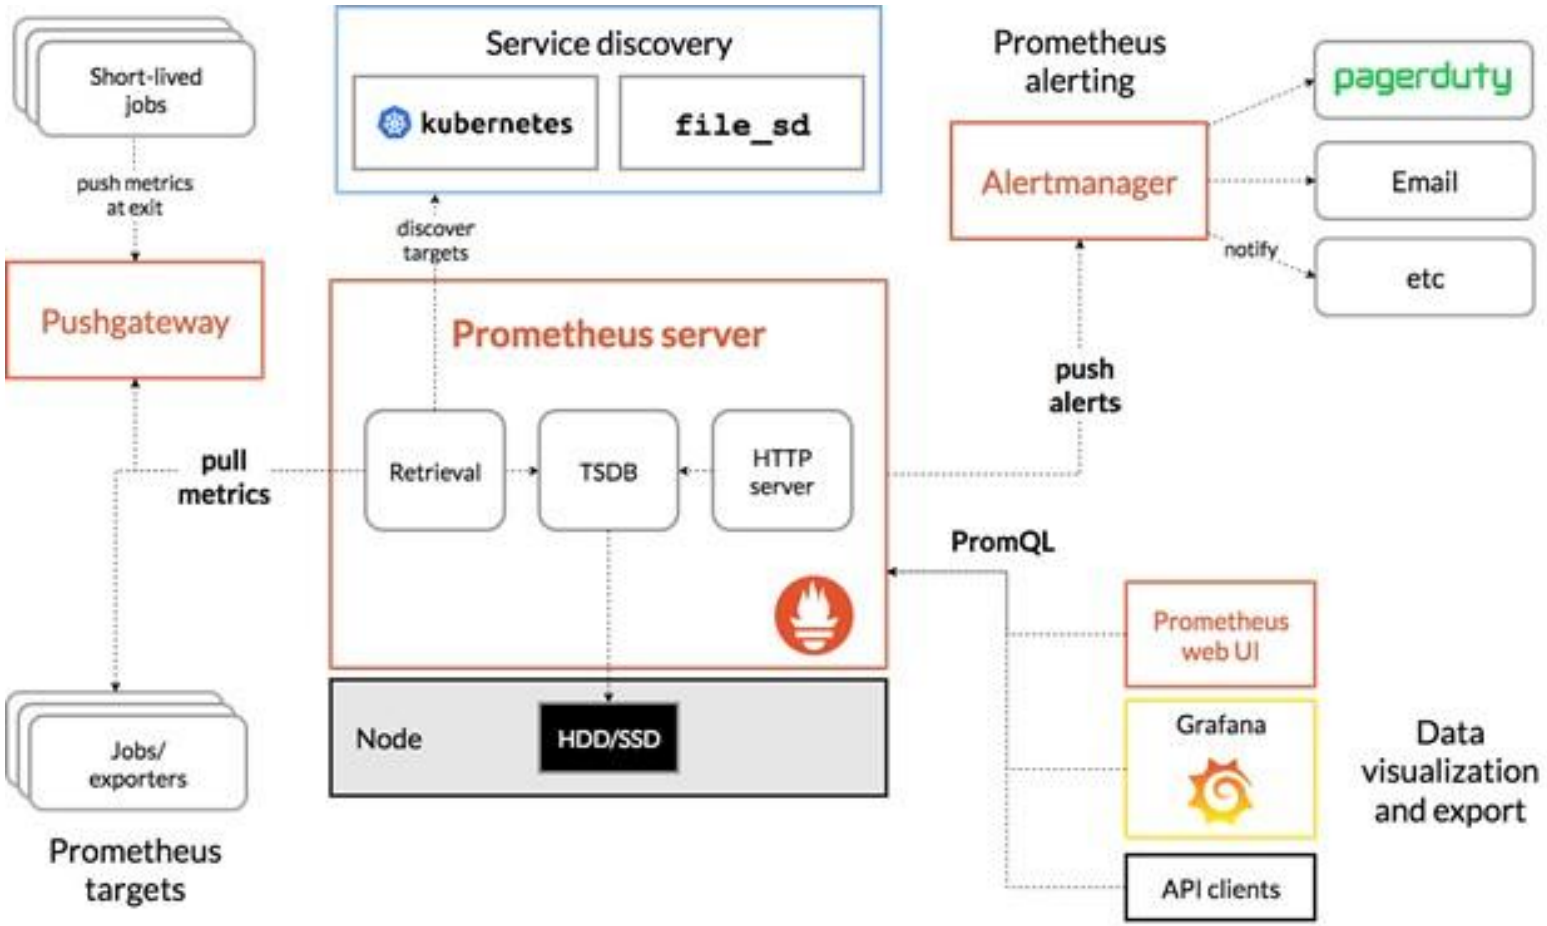

Deploy Prometheus and Grafana on Kubernetes

1. Creating all configuration YAML files and execute them in right order because of dependences

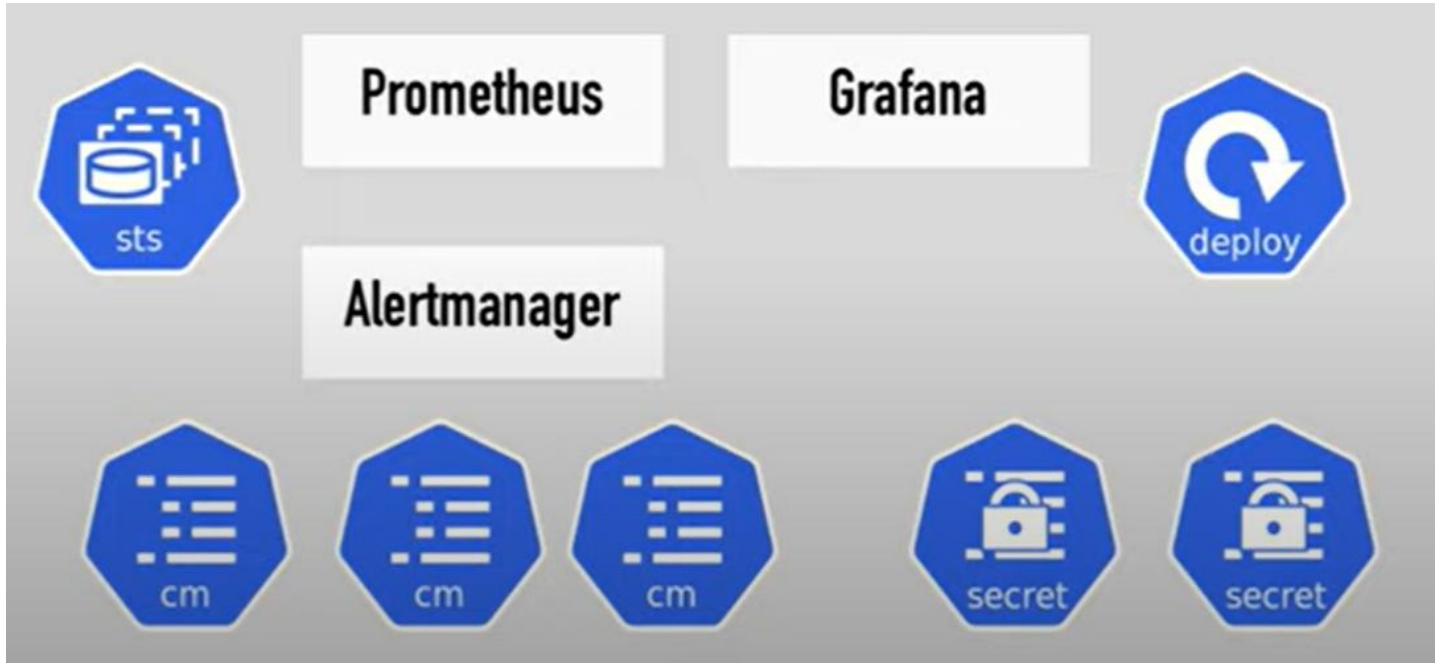

Deploy Prometheus and Grafana on Kubernetes

1. Creating all configuration YAML files and execute them in right order because of dependences

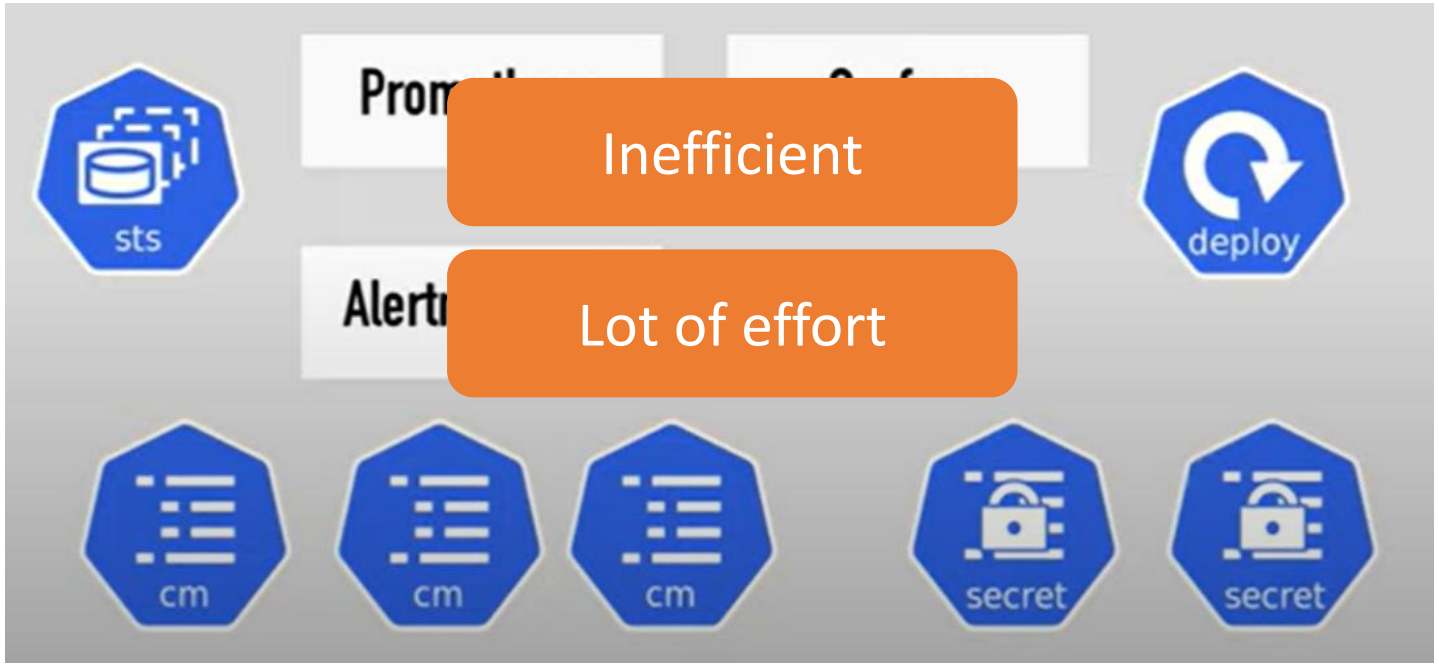

Deploy Prometheus and Grafana on Kubernetes

2. Using operators
   - Manages all Prometheus component
   - Consider all the Prometheus stack as one unit
   - You do not need to manually manage all these components
   - Find Prometheus operator
   - Deploy in the k8s cluster using the configuration files of the operator

Deploy Prometheus and Grafana on Kubernetes

3. Using helm chart to deploy the operator
   - Helm is popularly known as the package manager for Kubernetes.
   - Helm helps to configure and install any application to Kubernetes.
   - It bundles a Kubernetes application into a package known as a Helm Chart.
   - Helm Charts helps in defining, installing, and upgrading the most complex Kubernetes application
   - Go or Ansible scripts are an alternative of Helm
   - Helm: initiate setup
   - Operator: manage setup

Using helm chart to Deploy Prometheus

- Install Helm on Linux

```
snap install helm --classic
```

- Chart Folder contains the chart dependencies for the downloaded Helm Chart.
- Chart.yaml contains the meta-information about the downloaded Helm Chart such as name, description version number, and list of dependencies.
- Template Folder contains all the configuration YAML files you are deploying to the Kubernetes Cluster with the Helm Chart. These files are the Deployments, Services, ConfigMaps, and Secrets configurations YAML files.
- Values.yaml configure and set all the values for the configurations of YAML files. The configuration YAML files read their values from this file.

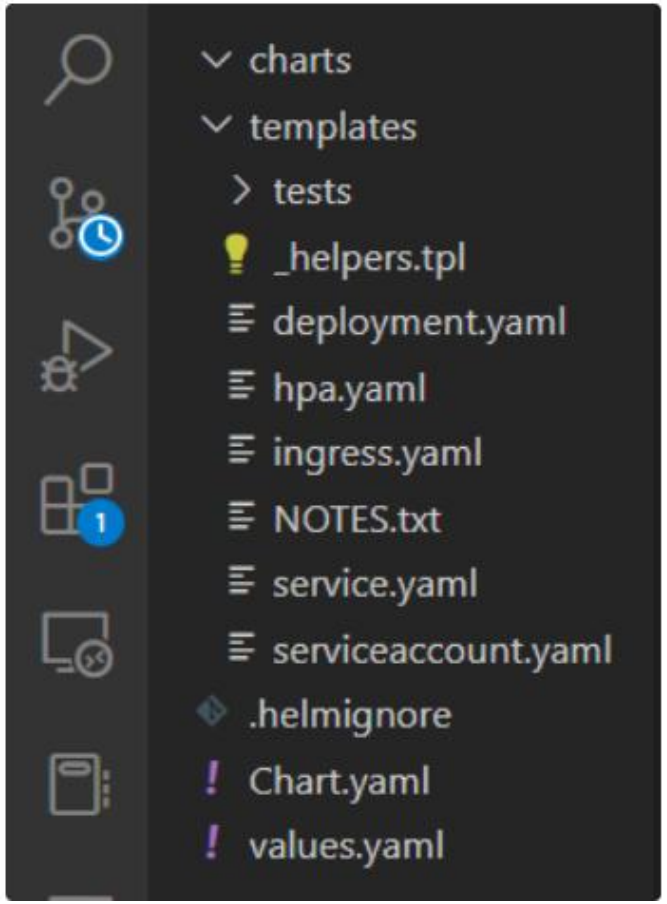

Using helm chart to Deploy Prometheus

- Update the helm repository

```
helm repo update
```

- Add Repositories. Add the following helm repositories by the commands below.

```
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo add stable https://charts.helm.sh/stable
```

- Update Helm Repositories

```
helm repo update
```

Using helm chart to Deploy Prometheus

- Install Prometheus Kubernetes

```
helm install prometheus prometheus-community/kube-prometheus-stack
```

- Get all the deployed Kubernetes resources for the Prometheus Kubernetes application

```
kubectl get all
```

Using helm chart to Deploy Prometheus

```
NAME                                                     READY   STATUS    RESTARTS   AGE
pod/alertmanager-prometheus-prometheus-oper-alertmanager-0   2/2     Running   0          16m
pod/prometheus-grafana-67596ff846-p8t6s                  2/2     Running   0          16m
pod/prometheus-kube-state-metrics-c65b87574-kprbb        1/1     Running   0          16m
pod/prometheus-prometheus-node-exporter-jr4mr            1/1     Running   0          16m
pod/prometheus-prometheus-oper-operator-7894f9c947-frn2x   2/2     Running   0          16m
pod/prometheus-prometheus-prometheus-oper-prometheus-0   3/3     Running   1          16m

NAME                                                TYPE        CLUSTER-IP       EXTERNAL-IP   PORT(S)                      AGE
service/alertmanager-operated                       ClusterIP   None             <none>        9093/TCP,9094/TCP,9094/UDP   16m
service/kubernetes                                  ClusterIP   10.96.0.1        <none>        443/TCP                      45d
service/prometheus-grafana                          ClusterIP   10.101.32.54     <none>        80/TCP                       16m
service/prometheus-kube-state-metrics               ClusterIP   10.102.225.147   <none>        8080/TCP                     16m
service/prometheus-operated                         ClusterIP   None             <none>        9090/TCP                     16m
service/prometheus-prometheus-node-exporter         ClusterIP   10.98.219.26     <none>        9100/TCP                     16m
service/prometheus-prometheus-oper-alertmanager     ClusterIP   10.106.31.202    <none>        9093/TCP                     16m
service/prometheus-prometheus-oper-operator         ClusterIP   10.99.122.30     <none>        8080/TCP,443/TCP             16m
service/prometheus-prometheus-oper-prometheus       ClusterIP   10.105.14.189    <none>        9090/TCP                     16m

NAME                                                    DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
daemonset.apps/prometheus-prometheus-node-exporter      1         1         1       1            1           <none>          16m

NAME                                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/prometheus-grafana                  1/1     1            1           16m
deployment.apps/prometheus-kube-state-metrics       1/1     1            1           16m
deployment.apps/prometheus-prometheus-oper-operator   1/1     1            1           16m

NAME                                                       DESIRED   CURRENT   READY   AGE
replicaset.apps/prometheus-grafana-67596ff846              1         1         1       16m
replicaset.apps/prometheus-kube-state-metrics-c65b87574    1         1         1       16m
replicaset.apps/prometheus-prometheus-oper-operator-7894f9c947   1         1         1       16m

NAME                                                          READY   AGE
statefulset.apps/alertmanager-prometheus-prometheus-oper-alertmanager   1/1     16m
statefulset.apps/prometheus-prometheus-prometheus-oper-prometheus       1/1     16m
```

Using helm chart to Deploy Prometheus

Using helm chart to Deploy Prometheus

Statefulsets

1. Core Prometheus server managed by operator
2. Alertmanager

```
NAME                                                     READY   STATUS    RESTARTS   AGE
pod/alertmanager-prometheus-prometheus-oper-alertmanager-0   2/2     Running   0          16m
pod/prometheus-grafana-67596ff846-p8t6s                  2/2     Running   0          16m
pod/prometheus-kube-state-metrics-c65b87574-kprbb        1/1     Running   0          16m
pod/prometheus-prometheus-node-exporter-jr4mr            1/1     Running   0          16m
pod/prometheus-prometheus-oper-operator-7894f9c947-frn2x 2/2     Running   0          16m
pod/prometheus-prometheus-prometheus-oper-prometheus-0   3/3     Running   1          16m

NAME                                                TYPE        CLUSTER-IP       EXTERNAL-IP   PORT(S)                      AGE
service/alertmanager-operated                       ClusterIP   None             <none>        9093/TCP,9094/TCP,9094/UDP   16m
service/kubernetes                                  ClusterIP   10.96.0.1        <none>        443/TCP                      45d
service/prometheus-grafana                          ClusterIP   10.101.32.54     <none>        80/TCP                       16m
service/prometheus-kube-state-metrics               ClusterIP   10.102.225.147   <none>        8080/TCP                     16m
service/prometheus-operated                         ClusterIP   None             <none>        9090/TCP                     16m
service/prometheus-prometheus-node-exporter         ClusterIP   10.98.219.26     <none>        9100/TCP                     16m
service/prometheus-prometheus-oper-alertmanager     ClusterIP   10.106.31.202    <none>        9093/TCP                     16m
service/prometheus-prometheus-oper-operator         ClusterIP   10.99.122.30     <none>        8080/TCP,443/TCP             16m
service/prometheus-prometheus-oper-prometheus       ClusterIP   10.105.14.189    <none>        9090/TCP                     16m

NAME                                                DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
daemonset.apps/prometheus-prometheus-node-exporter  1         1         1       1            1           <none>          16m

NAME                                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/prometheus-grafana                  1/1     1            1           16m
deployment.apps/prometheus-kube-state-metrics       1/1     1            1           16m
deployment.apps/prometheus-prometheus-oper-operator 1/1     1            1           16m

NAME                                                      DESIRED   CURRENT   READY   AGE
replicaset.apps/prometheus-grafana-67596ff846             1         1         1       16m
replicaset.apps/prometheus-kube-state-metrics-c65b87574   1         1         1       16m
replicaset.apps/prometheus-prometheus-oper-operator-7894f9c947   1         1         1       16m

NAME                                                          READY   AGE
statefulset.apps/alertmanager-prometheus-prometheus-oper-alertmanager   1/1     16m
statefulset.apps/prometheus-prometheus-prometheus-oper-prometheus       1/1     16m
```

Using helm chart to Deploy Prometheus

Statefulsets

1. Core Prometheus server managed by operator
2. Alertmanager

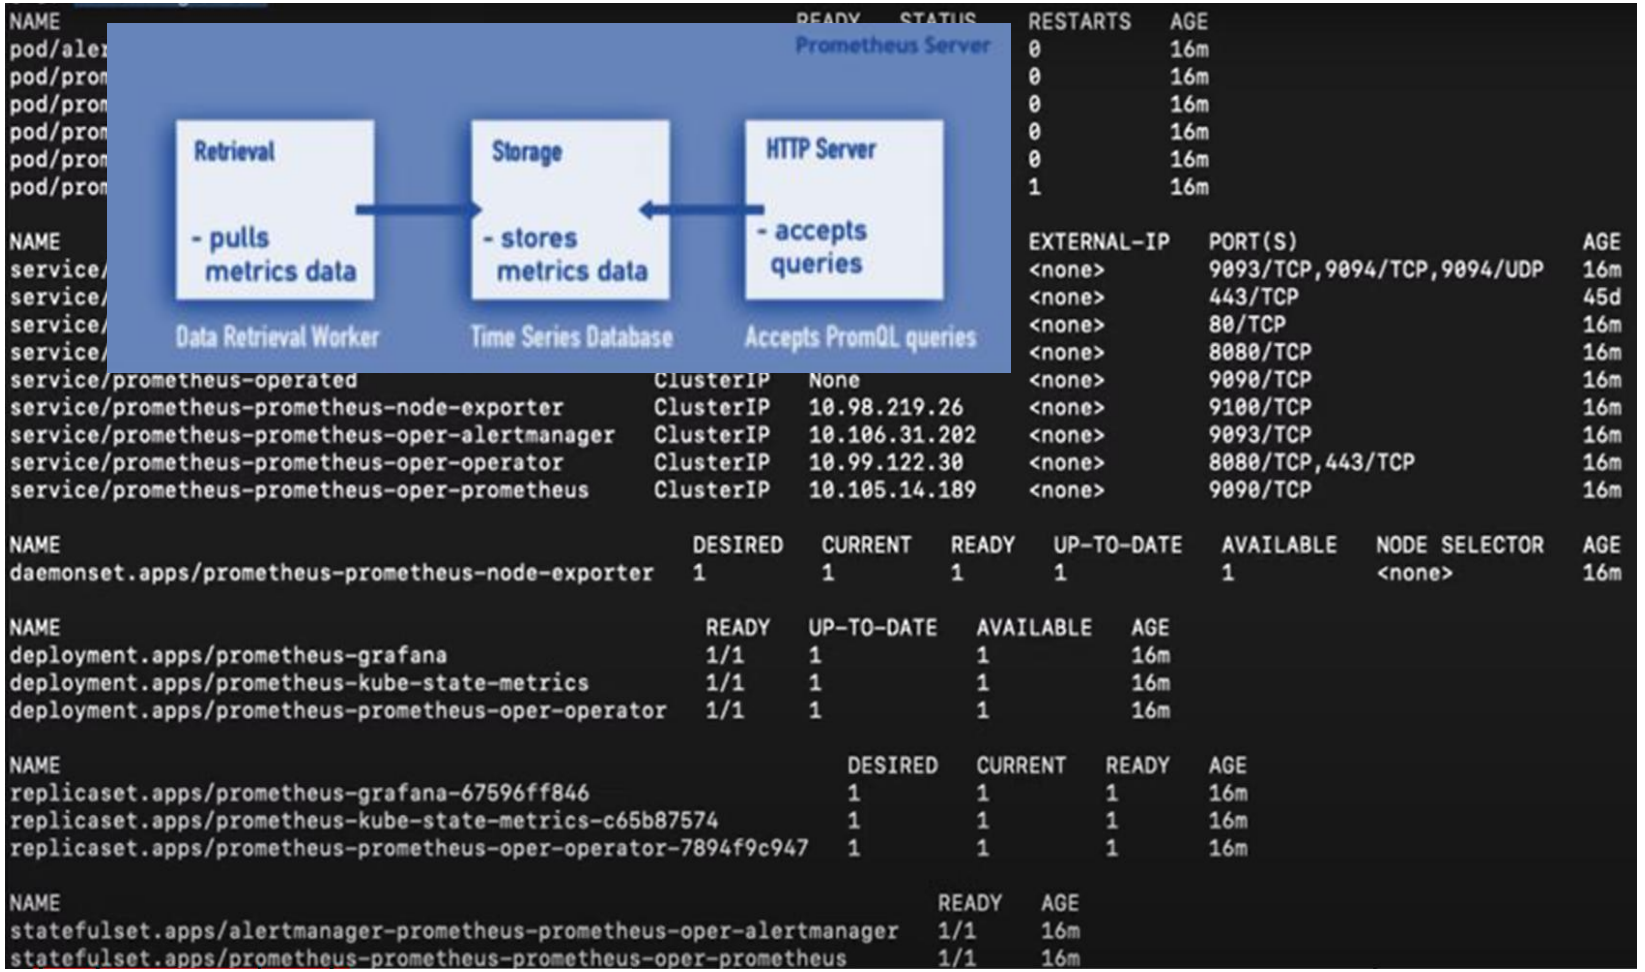

Using helm chart to Deploy Prometheus

Deployments

1. Prometheus operator
   - Created Prometheus and Alertmanager statefulset
2. Grafana
3. Kube state metrics
   - Own helm chart
   - Dependency of helm chart
   - Monitor the health of K8s component

```
zahra@zahra-laptop:~/.kube$ kubectl get all
NAME                                                     READY   STATUS    RESTARTS        AGE
pod/alertmanager-prometheus-kube-prometheus-alertmanager-0   2/2     Running   2 (7m58s ago)   20h
pod/mongodb-deployment-59d656b969-rcxrw                  1/1     Running   3 (7m58s ago)   3d3h
pod/nginx-deployment-cbdccf466-dnppg                     1/1     Running   6 (7m58s ago)   6d6h
pod/prometheus-grafana-6fdbffb5c9-5xzps                  3/3     Running   3 (7m58s ago)   20h
pod/prometheus-kube-prometheus-operator-757f8788d4-vms7c     1/1     Running   1 (7m58s ago)   20h
pod/prometheus-kube-state-metrics-898dd9b88-x6zqs        1/1     Running   1 (7m58s ago)   20h
pod/prometheus-prometheus-kube-prometheus-prometheus-0   2/2     Running   2 (7m58s ago)   20h
pod/prometheus-prometheus-node-exporter-bmgwt            1/1     Running   1 (7m58s ago)   20h

NAME                                              TYPE           CLUSTER-IP       EXTERNAL-IP   PORT(S)                      AGE
service/alertmanager-operated                     ClusterIP      None             <none>        9093/TCP,9094/TCP,9094/UDP   20h
service/kubernetes                                ClusterIP      10.96.0.1        <none>        443/TCP                      6d21h
service/mongo-experss-service                     LoadBalancer   10.97.175.65     <pending>     8081:30000/TCP               5d22h
service/mongodb-service                           ClusterIP      10.101.132.177   <none>        27017/TCP                    6d1h
service/mongoex-service                           ClusterIP      10.99.154.199    <none>        8081/TCP                     3d19h
service/nginx-service                             ClusterIP      10.107.137.84    <none>        80/TCP                       6d5h
service/prometheus-grafana                        ClusterIP      10.106.179.180   <none>        80/TCP                       20h
service/prometheus-kube-prometheus-alertmanager   ClusterIP      10.109.215.82    <none>        9093/TCP,8080/TCP            20h
service/prometheus-kube-prometheus-operator       ClusterIP      10.99.28.217     <none>        443/TCP                      20h
service/prometheus-kube-prometheus-prometheus     ClusterIP      10.99.105.32     <none>        9090/TCP,8080/TCP            20h
service/prometheus-kube-state-metrics             ClusterIP      10.102.186.174   <none>        8080/TCP                     20h
service/prometheus-operated                       ClusterIP      None             <none>        9090/TCP                     20h
service/prometheus-prometheus-node-exporter       ClusterIP      10.103.180.36    <none>        9100/TCP                     20h

NAME                                                 DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR            AGE
daemonset.apps/prometheus-prometheus-node-exporter   1         1         1       1            1           kubernetes.io/os=linux   20h

NAME                                                 READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mongodb-deployment                   1/1     1            1           3d17h
deployment.apps/nginx-deployment                     1/1     1            1           6d6h
deployment.apps/prometheus-grafana                   1/1     1            1           20h
deployment.apps/prometheus-kube-prometheus-operator  1/1     1            1           20h
deployment.apps/prometheus-kube-state-metrics        1/1     1            1           20h

NAME                                                       DESIRED   CURRENT   READY   AGE
replicaset.apps/mongodb-deployment-59d656b969              1         1         1       3d3h
replicaset.apps/mongodb-deployment-766d54c95               0         0         0       3d17h
replicaset.apps/nginx-deployment-cbdccf466                 1         1         1       6d6h
replicaset.apps/prometheus-grafana-6fdbffb5c9              1         1         1       20h
replicaset.apps/prometheus-kube-prometheus-operator-757f8788d4   1         1         1       20h
replicaset.apps/prometheus-kube-state-metrics-898dd9b88    1         1         1       20h

NAME                                                          READY   AGE
statefulset.apps/alertmanager-prometheus-kube-prometheus-alertmanager   1/1     20h
statefulset.apps/prometheus-prometheus-kube-prometheus-prometheus       1/1     20h
```

Using helm chart to Deploy Prometheus

- Replicaset
  - Created by deployment

```
zahra@zahra-laptop:~/.kube$ kubectl get all
NAME                                                     READY   STATUS    RESTARTS        AGE
pod/alertmanager-prometheus-kube-prometheus-alertmanager-0   2/2     Running   2 (7m58s ago)   20h
pod/mongodb-deployment-59d656b969-rcxrw                  1/1     Running   3 (7m58s ago)   3d3h
pod/nginx-deployment-cbdccf466-dnppg                     1/1     Running   6 (7m58s ago)   6d6h
pod/prometheus-grafana-6fdbffb5c9-5xzps                  3/3     Running   3 (7m58s ago)   20h
pod/prometheus-kube-prometheus-operator-757f8788d4-vms7c     1/1     Running   1 (7m58s ago)   20h
pod/prometheus-kube-state-metrics-898dd9b88-x6zqs        1/1     Running   1 (7m58s ago)   20h
pod/prometheus-prometheus-kube-prometheus-prometheus-0   2/2     Running   2 (7m58s ago)   20h
pod/prometheus-prometheus-node-exporter-bmgwt            1/1     Running   1 (7m58s ago)   20h

NAME                                              TYPE           CLUSTER-IP       EXTERNAL-IP   PORT(S)                      AGE
service/alertmanager-operated                     ClusterIP      None             <none>        9093/TCP,9094/TCP,9094/UDP   20h
service/kubernetes                                ClusterIP      10.96.0.1        <none>        443/TCP                      6d21h
service/mongo-experss-service                     LoadBalancer   10.97.175.65     <pending>     8081:30000/TCP               5d22h
service/mongodb-service                           ClusterIP      10.101.132.177   <none>        27017/TCP                    6d1h
service/mongoex-service                           ClusterIP      10.99.154.199    <none>        8081/TCP                     3d19h
service/nginx-service                             ClusterIP      10.107.137.84    <none>        80/TCP                       6d5h
service/prometheus-grafana                        ClusterIP      10.106.179.180   <none>        80/TCP                       20h
service/prometheus-kube-prometheus-alertmanager   ClusterIP      10.109.215.82    <none>        9093/TCP,8080/TCP            20h
service/prometheus-kube-prometheus-operator       ClusterIP      10.99.28.217     <none>        443/TCP                      20h
service/prometheus-kube-prometheus-prometheus     ClusterIP      10.99.105.32     <none>        9090/TCP,8080/TCP            20h
service/prometheus-kube-state-metrics             ClusterIP      10.102.186.174   <none>        8080/TCP                     20h
service/prometheus-operated                       ClusterIP      None             <none>        9090/TCP                     20h
service/prometheus-prometheus-node-exporter       ClusterIP      10.103.180.36    <none>        9100/TCP                     20h

NAME                                                 DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR            AGE
daemonset.apps/prometheus-prometheus-node-exporter   1         1         1       1            1           kubernetes.io/os=linux   20h

NAME                                                 READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mongodb-deployment                   1/1     1            1           3d17h
deployment.apps/nginx-deployment                     1/1     1            1           6d6h
deployment.apps/prometheus-grafana                   1/1     1            1           20h
deployment.apps/prometheus-kube-prometheus-operator  1/1     1            1           20h
deployment.apps/prometheus-kube-state-metrics        1/1     1            1           20h

NAME                                                           DESIRED   CURRENT   READY   AGE
replicaset.apps/mongodb-deployment-59d656b969                  1         1         1       3d3h
replicaset.apps/mongodb-deployment-766d54c95                   0         0         0       3d17h
replicaset.apps/nginx-deployment-cbdccf466                     1         1         1       6d6h
replicaset.apps/prometheus-grafana-6fdbffb5c9                  1         1         1       20h
replicaset.apps/prometheus-kube-prometheus-operator-757f8788d4   1         1         1       20h
replicaset.apps/prometheus-kube-state-metrics-898dd9b88        1         1         1       20h

NAME                                                              READY   AGE
statefulset.apps/alertmanager-prometheus-kube-prometheus-alertmanager   1/1     20h
statefulset.apps/prometheus-prometheus-kube-prometheus-prometheus       1/1     20h
```

Using helm chart to Deploy Prometheus

Daemonset

- Node exporter daemonset runs on every node
- Connected to the server
- Translates worker node metrics(CPU usage, load on server) to Prometheus metrics

```
zahra@zahra-laptop:~/.kube$ kubectl get all
NAME                                                     READY   STATUS    RESTARTS        AGE
pod/alertmanager-prometheus-kube-prometheus-alertmanager-0   2/2     Running   2 (7m58s ago)   20h
pod/mongodb-deployment-59d656b969-rcxrw                  1/1     Running   3 (7m58s ago)   3d3h
pod/nginx-deployment-cbdccf466-dnppg                     1/1     Running   6 (7m58s ago)   6d6h
pod/prometheus-grafana-6fdbffb5c9-5xzps                  3/3     Running   3 (7m58s ago)   20h
pod/prometheus-kube-prometheus-operator-757f8788d4-vms7c     1/1     Running   1 (7m58s ago)   20h
pod/prometheus-kube-state-metrics-898dd9b88-x6zqs        1/1     Running   1 (7m58s ago)   20h
pod/prometheus-prometheus-kube-prometheus-prometheus-0   2/2     Running   2 (7m58s ago)   20h
pod/prometheus-prometheus-node-exporter-bmgwt            1/1     Running   1 (7m58s ago)   20h

NAME                                              TYPE           CLUSTER-IP       EXTERNAL-IP   PORT(S)                      AGE
service/alertmanager-operated                     ClusterIP      None             <none>        9093/TCP,9094/TCP,9094/UDP   20h
service/kubernetes                                ClusterIP      10.96.0.1        <none>        443/TCP                      6d21h
service/mongo-experss-service                     LoadBalancer   10.97.175.65     <pending>     8081:30000/TCP               5d22h
service/mongodb-service                           ClusterIP      10.101.132.177   <none>        27017/TCP                    6d1h
service/mongoex-service                           ClusterIP      10.99.154.199    <none>        8081/TCP                     3d19h
service/nginx-service                             ClusterIP      10.107.137.84    <none>        80/TCP                       6d5h
service/prometheus-grafana                        ClusterIP      10.106.179.180   <none>        80/TCP                       20h
service/prometheus-kube-prometheus-alertmanager   ClusterIP      10.109.215.82    <none>        9093/TCP,8080/TCP            20h
service/prometheus-kube-prometheus-operator       ClusterIP      10.99.28.217     <none>        443/TCP                      20h
service/prometheus-kube-prometheus-prometheus     ClusterIP      10.99.105.32     <none>        9090/TCP,8080/TCP            20h
service/prometheus-kube-state-metrics             ClusterIP      10.102.186.174   <none>        8080/TCP                     20h
service/prometheus-operated                       ClusterIP      None             <none>        9090/TCP                     20h
service/prometheus-prometheus-node-exporter       ClusterIP      10.103.180.36    <none>        9100/TCP                     20h

NAME                                                DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR            AGE
daemonset.apps/prometheus-prometheus-node-exporter   1         1         1       1            1           kubernetes.io/os=linux   20h

NAME                                             READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mongodb-deployment               1/1     1            1           3d17h
deployment.apps/nginx-deployment                 1/1     1            1           6d6h
deployment.apps/prometheus-grafana               1/1     1            1           20h
deployment.apps/prometheus-kube-prometheus-operator   1/1     1            1           20h
deployment.apps/prometheus-kube-state-metrics    1/1     1            1           20h

NAME                                                        DESIRED   CURRENT   READY   AGE
replicaset.apps/mongodb-deployment-59d656b969               1         1         1       3d3h
replicaset.apps/mongodb-deployment-766d54c95                0         0         0       3d17h
replicaset.apps/nginx-deployment-cbdccf466                  1         1         1       6d6h
replicaset.apps/prometheus-grafana-6fdbffb5c9               1         1         1       20h
replicaset.apps/prometheus-kube-prometheus-operator-757f8788d4   1         1         1       20h
replicaset.apps/prometheus-kube-state-metrics-898dd9b88     1         1         1       20h

NAME                                                          READY   AGE
statefulset.apps/alertmanager-prometheus-kube-prometheus-alertmanager   1/1     20h
statefulset.apps/prometheus-prometheus-kube-prometheus-prometheus       1/1     20h
```

Using helm chart to Deploy Prometheus

Configmap

- Configuration for different parts
  - Managing by operator
  - Containing the information for connecting to default metrics

```
zahra@zahra-laptop:~/.kube$ kubectl get configmap
NAME                                                  DATA   AGE
kube-root-ca.crt                                      1      6d21h
mongodb-configmap                                     1      6d
prometheus-grafana                                    1      20h
prometheus-grafana-config-dashboards                  1      20h
prometheus-kube-prometheus-alertmanager-overview      1      20h
prometheus-kube-prometheus-apiserver                  1      20h
prometheus-kube-prometheus-cluster-total              1      20h
prometheus-kube-prometheus-controller-manager         1      20h
prometheus-kube-prometheus-etcd                       1      20h
prometheus-kube-prometheus-grafana-datasource         1      20h
prometheus-kube-prometheus-grafana-overview           1      20h
prometheus-kube-prometheus-k8s-coredns                1      20h
prometheus-kube-prometheus-k8s-resources-cluster      1      20h
prometheus-kube-prometheus-k8s-resources-multicluster 1      20h
prometheus-kube-prometheus-k8s-resources-namespace    1      20h
prometheus-kube-prometheus-k8s-resources-node         1      20h
prometheus-kube-prometheus-k8s-resources-pod          1      20h
prometheus-kube-prometheus-k8s-resources-workload     1      20h
prometheus-kube-prometheus-k8s-resources-workloads-namespace 1 20h
prometheus-kube-prometheus-kubelet                    1      20h
prometheus-kube-prometheus-namespace-by-pod           1      20h
prometheus-kube-prometheus-namespace-by-workload      1      20h
prometheus-kube-prometheus-node-cluster-rsrc-use      1      20h
prometheus-kube-prometheus-node-rsrc-use              1      20h
prometheus-kube-prometheus-nodes                      1      20h
prometheus-kube-prometheus-nodes-darwin               1      20h
prometheus-kube-prometheus-persistentvolumesusage     1      20h
prometheus-kube-prometheus-pod-total                  1      20h
prometheus-kube-prometheus-prometheus                 1      20h
prometheus-kube-prometheus-proxy                      1      20h
prometheus-kube-prometheus-scheduler                  1      20h
prometheus-kube-prometheus-workload-total             1      20h
prometheus-prometheus-kube-prometheus-prometheus-rulefiles-0 29  20h
```

Using helm chart to Deploy Prometheus

Secrets

- For Grafana
- For Prometheus
- For operator

Including

- Certificates
- Username and passwords

```
zahra@zahra-laptop:~/.kube$ kubectl get secret
NAME                                                         TYPE                 DATA   AGE
alertmanager-prometheus-kube-prometheus-alertmanager         Opaque               1      20h
alertmanager-prometheus-kube-prometheus-alertmanager-generated   Opaque          1      20h
alertmanager-prometheus-kube-prometheus-alertmanager-tls-assets-0   Opaque       0      20h
alertmanager-prometheus-kube-prometheus-alertmanager-web-config   Opaque         1      20h
kube-prometheus-stack-admission                              Opaque               3      22h
mongo-express-secret                                         Opaque               2      6d1h
mongodb-secret                                               Opaque               2      3d17h
prometheus-grafana                                           Opaque               3      20h
prometheus-kube-prometheus-admission                         Opaque               3      20h
prometheus-kube-prometheus-prometheus                        Opaque               0      20h
prometheus-prometheus-kube-prometheus-prometheus             Opaque               1      20h
prometheus-prometheus-kube-prometheus-prometheus-tls-assets-0   Opaque            1      20h
prometheus-prometheus-kube-prometheus-prometheus-web-config   Opaque              1      20h
sh.helm.release.v1.prometheus.v1                             helm.sh/release.v1   1      20h
```

Using helm chart to Deploy Prometheus

What is inside

- Prometheus?
- Alertmanager?
- Operator?

Kubctl describe component-type component-name

```
[~]$ kubectl get statefulset
NAME                                                READY   AGE
alertmanager-prometheus-prometheus-oper-alertmanager   1/1     18h
prometheus-prometheus-prometheus-oper-prometheus       1/1     18h
[~]$ kubectl describe statefulset prometheus-prometheus-prometheus-oper-prometheus > prom.yaml
[~]$ kubectl describe statefulset alertmanager-prometheus-prometheus-oper-alertmanager > alert.yaml
[~]$ kubectl get deployment
NAME                                  READY   UP-TO-DATE   AVAILABLE   AGE
prometheus-grafana                    1/1     1            1           18h
prometheus-kube-state-metrics         1/1     1            1           18h
prometheus-prometheus-oper-operator   1/1     1            1           18h
[~]$ kubectl describe deployment prometheus-prometheus-oper-operator > oper.yaml
```

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

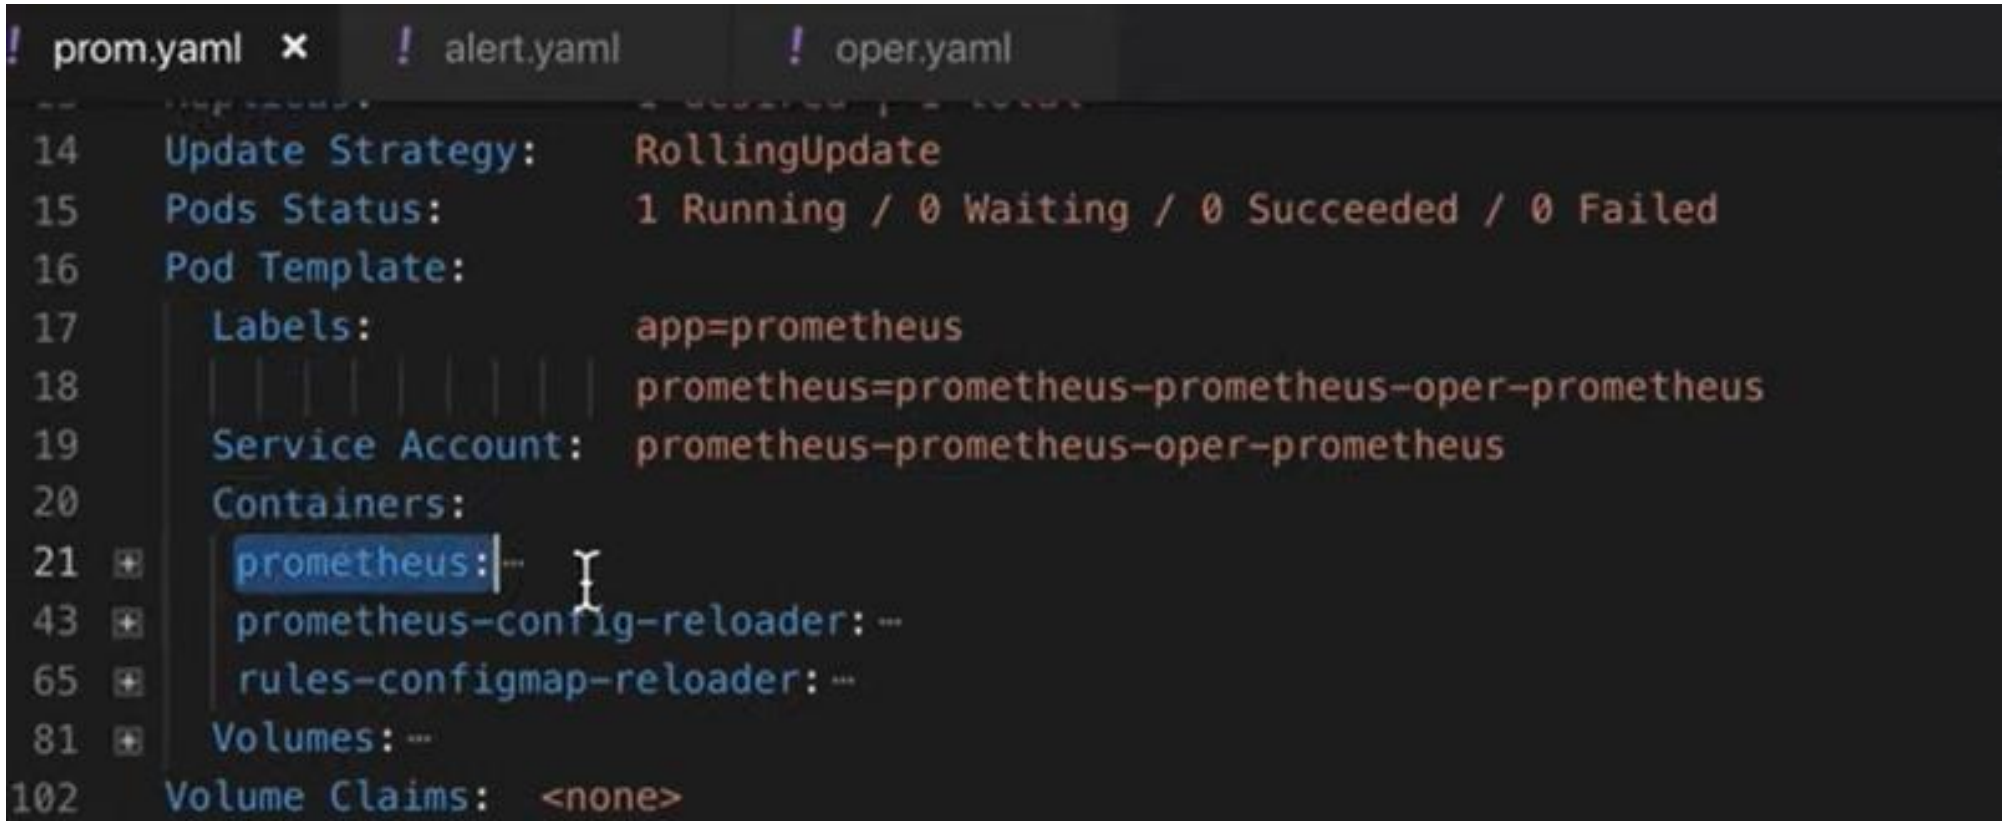

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

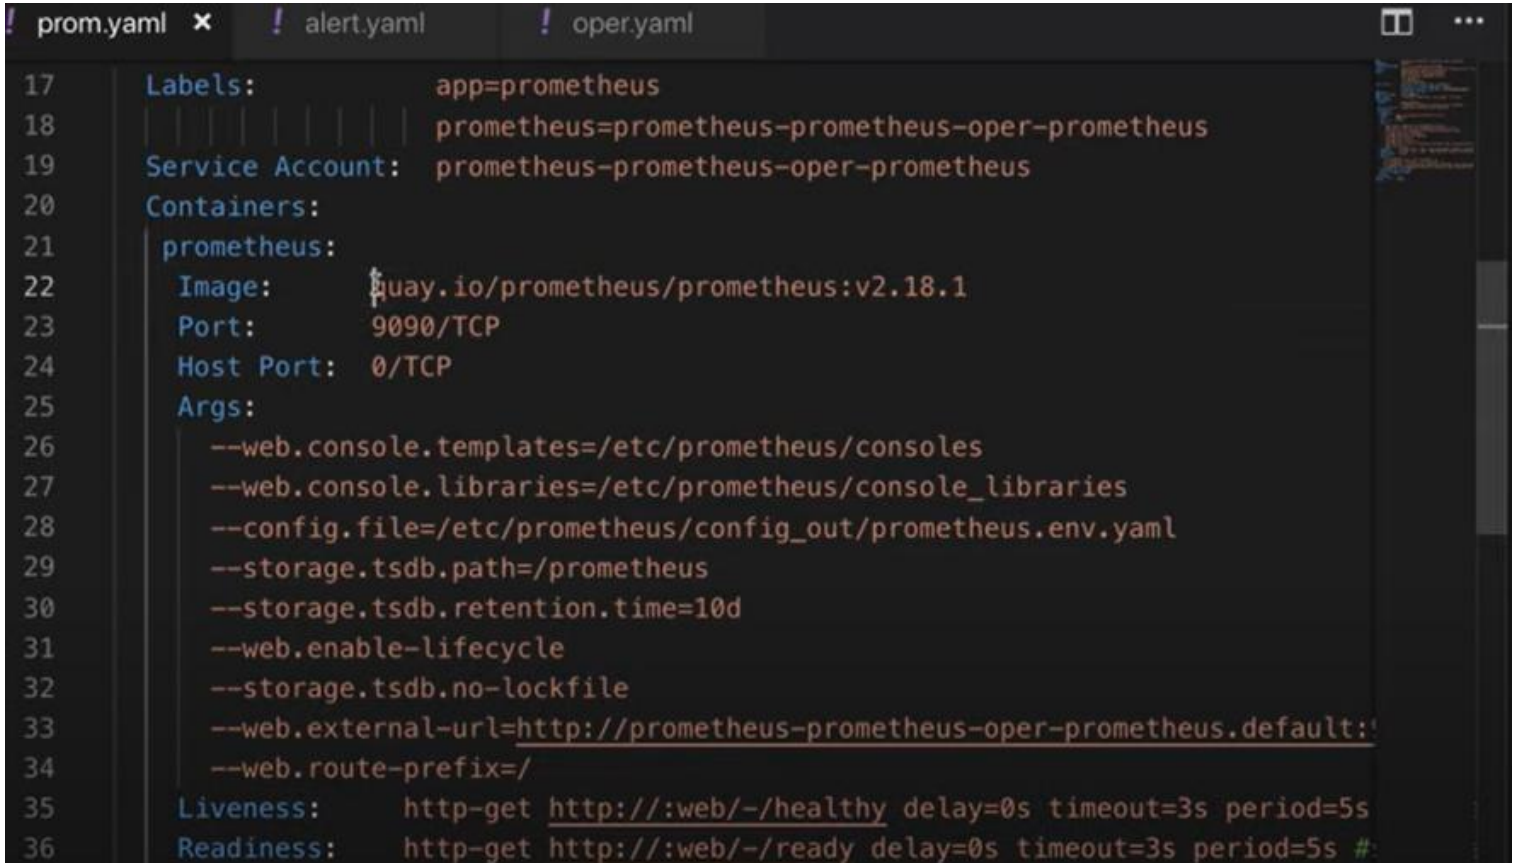

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

- Mount: Prometheus gets all the configuration data
  - Configuration file defines what end point to scrape. It has all the addresses of all the applications that they expose this metric endpoint.
  - Rules file define different rules (alerting rules, etc)

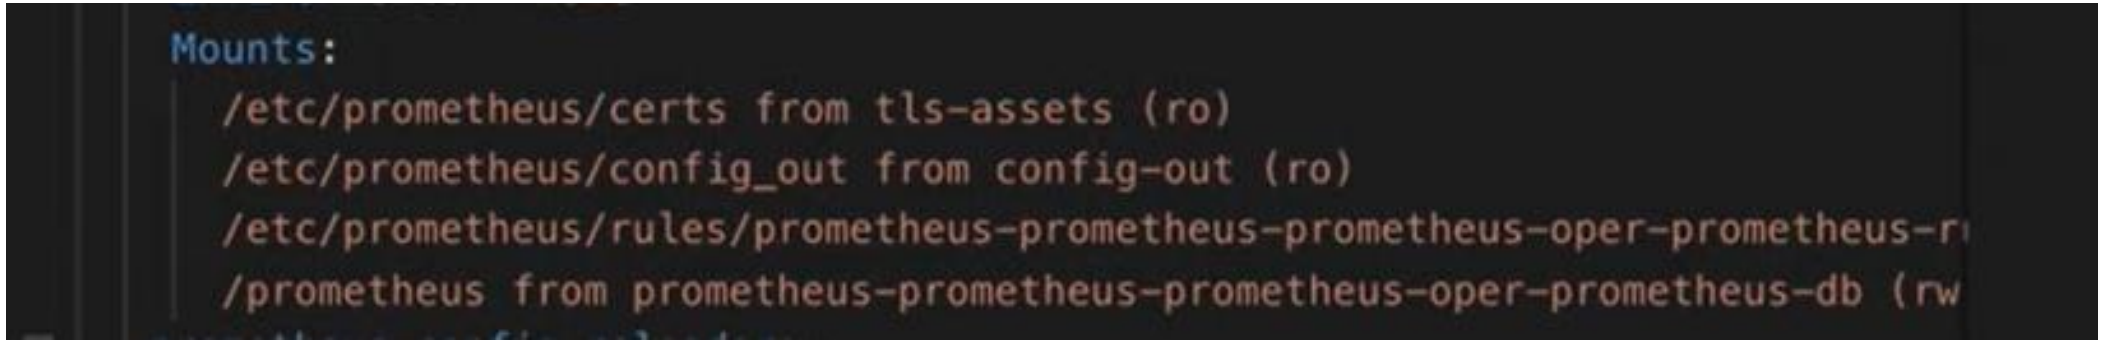

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

```
prometheus-config-reloader: --
rules-configmap-reloader:   --
```

They help to reload the main container of Prometheus when there is some changes in the configuration.

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

The rules-configmap-reloader: shows all the rules

Mount: shows the pass for the rule file

```
rules-configmap-reloader:
 Image:      docker.io/jimmidyson/configmap-reload:v0.3.0
 Port:       <none>
 Host Port:  <none>
 Args:
   --webhook-url=http://127.0.0.1:9090/-/reload
   --volume-dir=/etc/prometheus/rules/prometheus-prometheus-prometheus-oper-prometheus
 Limits:
   cpu:     100m
   memory:  25Mi
 Requests:
   cpu:        100m
   memory:     25Mi
 Environment:  <none>
 Mounts:
   /etc/prometheus/rules/prometheus-prometheus-prometheus-oper-prometheus-rulefiles-0
```

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

The rules-configmap-reloader: shows all the rules

Mount: shows the pass for the rule file

```
[~]$ kubectl get configmap prometheus-prometheus-prometheus-oper-prometheus-rulefiles-0 -o yaml > confg.yaml
[~]$
```

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

- rules file

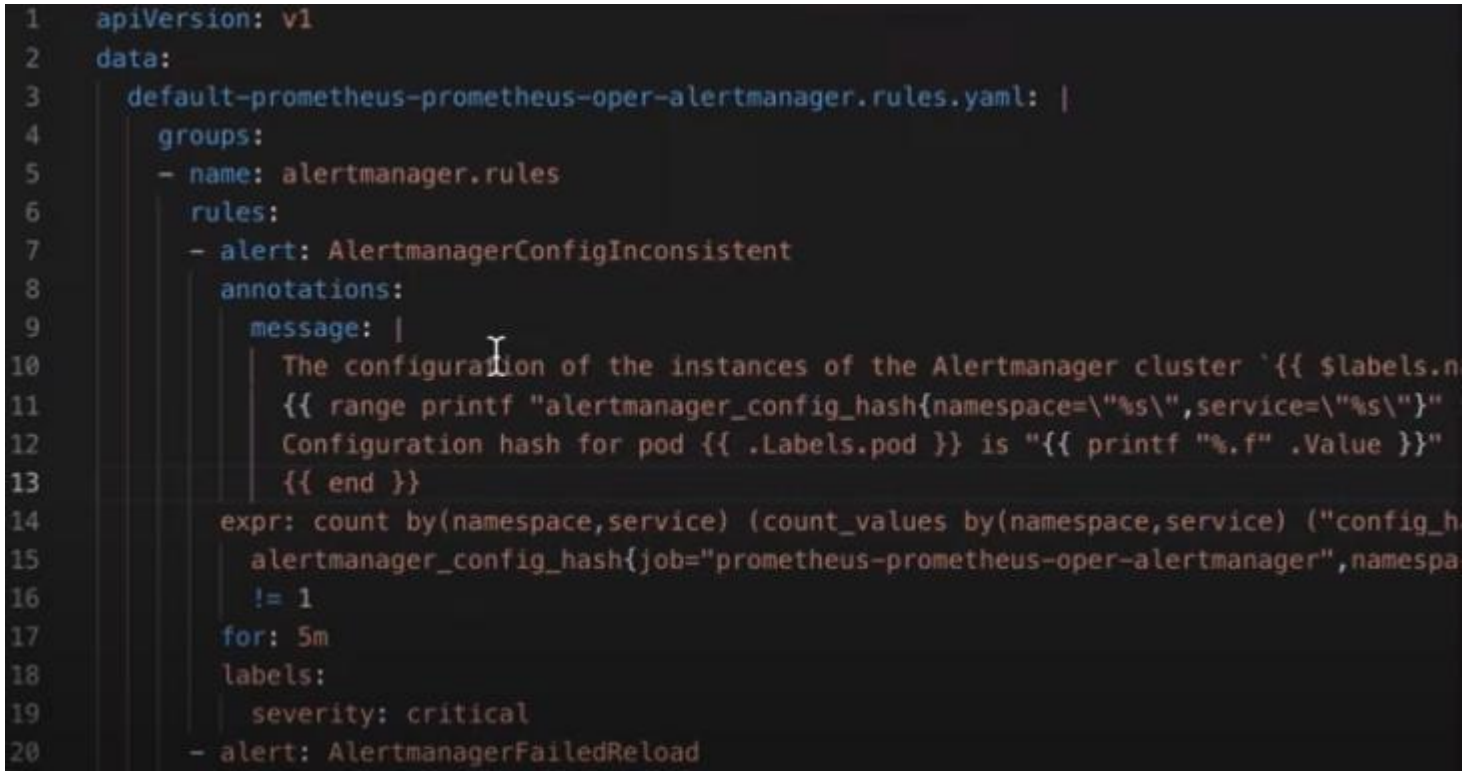

```
apiVersion: v1
data:
  default-prometheus-prometheus-oper-alertmanager.rules.yaml: |
    groups:
    - name: alertmanager.rules
      rules:
      - alert: AlertmanagerConfigInconsistent
        annotations:
          message: |
            The configuration of the instances of the Alertmanager cluster `{{ $labels.n
            {{ range printf "alertmanager_config_hash{namespace=\"%s\",service=\"%s\"}"
            Configuration hash for pod {{ .Labels.pod }} is "{{ printf "%.f" .Value }}"
            {{ end }}
        expr: count by(namespace,service) (count_values by(namespace,service) ("config_h
          alertmanager_config_hash{job="prometheus-prometheus-oper-alertmanager",namespa
          != 1
        for: 5m
        labels:
          severity: critical
      - alert: AlertmanagerFailedReload
```

Using helm chart to Deploy Prometheus

What is inside Prometheus Yaml file?

- Alert manager

```
Name:                   alertmanager-prometheus-prometheus-oper-alertmanager
Namespace:              default
CreationTimestamp:      Fri, 17 Jul 2020 16:11:59 +0200
Selector:               alertmanager=prometheus-prometheus-oper-alertmanager,app=alertmanager
Labels:                 app=prometheus-operator-alertmanager
                        app.kubernetes.io/managed-by=Helm
                        chart=prometheus-operator-8.15.12
                        heritage=Helm
                        release=prometheus
Annotations:            meta.helm.sh/release-name: prometheus
                        meta.helm.sh/release-namespace: default
Replicas:               1 desired | 1 total
Update Strategy:        RollingUpdate
Pods Status:            1 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:               alertmanager=prometheus-prometheus-oper-alertmanager
                        app=alertmanager
  Service Account:      prometheus-prometheus-oper-alertmanager
  Containers:
   alertmanager:
```

Accessing Grafana

Grafana is UI providing data visualization for Prometheus data

```
[~]$ kubectl get service
NAME                                      TYPE        CLUSTER-IP       EXTERNAL-IP   PORT(S)
alertmanager-operated                     ClusterIP   None             <none>        9093/TCP,9094/TCP,9094/UDP
kubernetes                                ClusterIP   10.96.0.1        <none>        443/TCP
prometheus-grafana                        ClusterIP   10.101.32.54     <none>        80/TCP
prometheus-kube-state-metrics             ClusterIP   10.102.225.147   <none>        8080/TCP
prometheus-operated                       ClusterIP   None             <none>        9090/TCP
prometheus-prometheus-node-exporter       ClusterIP   10.98.219.26     <none>        9100/TCP
prometheus-prometheus-oper-alertmanager   ClusterIP   10.106.31.202    <none>        9093/TCP
prometheus-prometheus-oper-operator       ClusterIP   10.99.122.30     <none>        8080/TCP,443/TCP
prometheus-prometheus-oper-prometheus     ClusterIP   10.105.14.189    <none>        9090/TCP
[~]$
```

Internal service

Accessing Grafana

Access Grafana through port forward:

1. Get Grafana pod logs

```
[~]$ kubectl get pod
NAME                                                     READY   STATUS    RESTARTS   AGE
alertmanager-prometheus-prometheus-oper-alertmanager-0   2/2     Running   0          18h
prometheus-grafana-67596ff846-p8t6s                      2/2     Running   0          18h
prometheus-kube-state-metrics-c65b87574-kprbb            1/1     Running   0          18h
prometheus-prometheus-node-exporter-jr4mr                1/1     Running   0          18h
prometheus-prometheus-oper-operator-7894f9c947-frn2x     2/2     Running   0          18h
prometheus-prometheus-prometheus-oper-prometheus-0       3/3     Running   1          18h
[~]$ kubectl logs prometheus-grafana-67596ff846-p8t6s
error: a container name must be specified for pod prometheus-grafana-67596ff846-p8t6s, choose one of: [grafana
oard grafana] or one of the init containers: [grafana-sc-datasources]
[~]$ kubectl logs prometheus-grafana-67596ff846-p8t6s -c grafan
```

Accessing Grafana

Getting Grafana port and username

- Username: admin
- password: prom-operator(https://github.com/helm/charts/tree/master/stable/prometheus-operator)

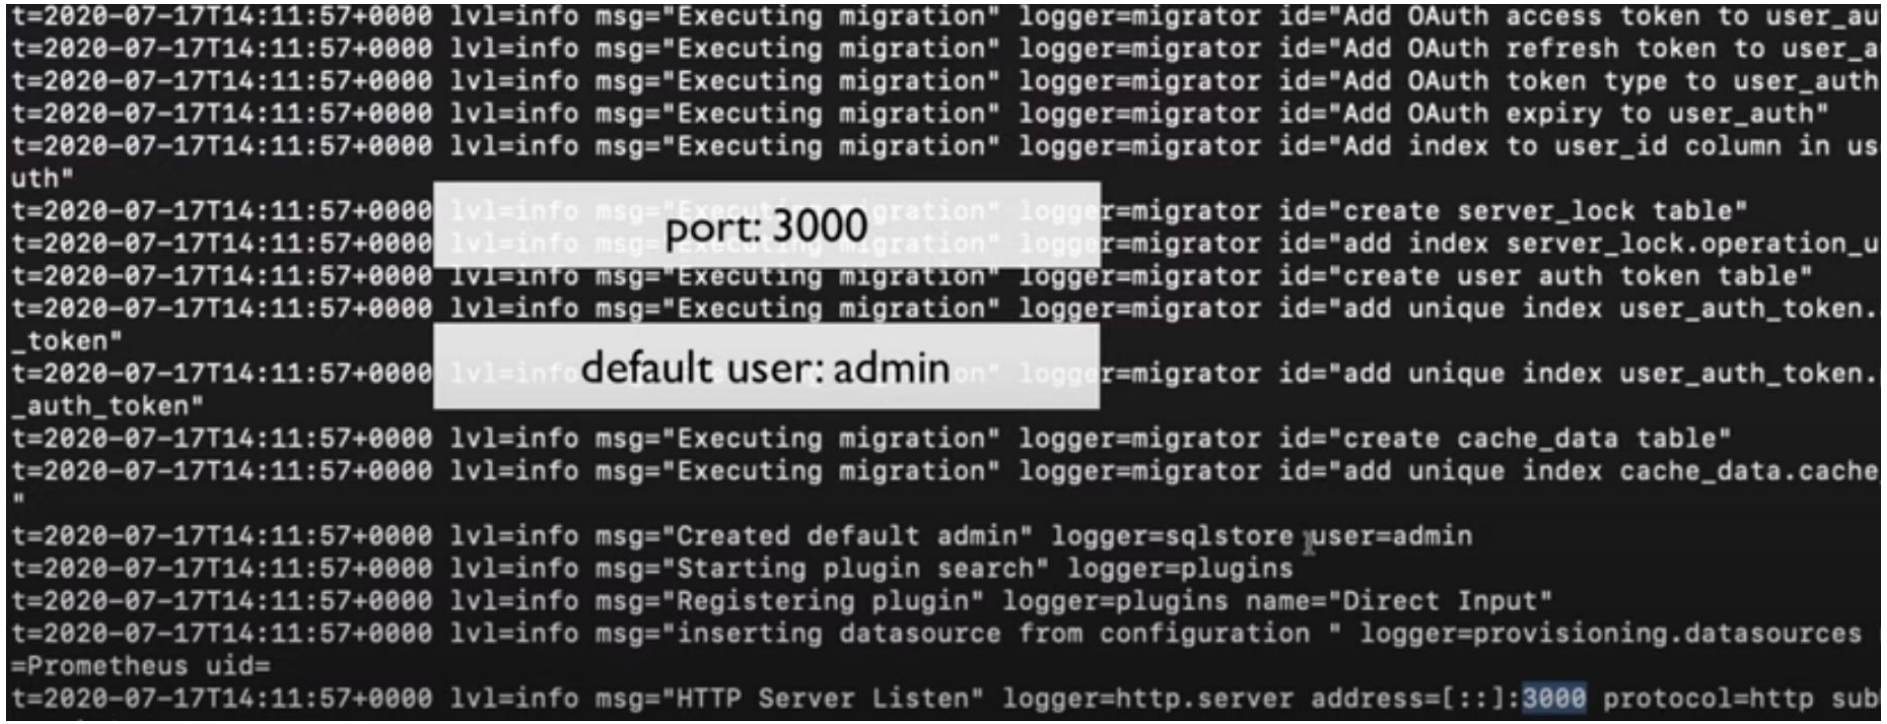

Accessing Grafana

Opening the port 3000

```
[~]$ kubectl port-forward deployment/prometheus-grafana 3000
Forwarding from 127.0.0.1:3000 -> 3000
Forwarding from [::1]:3000 -> 3000
```

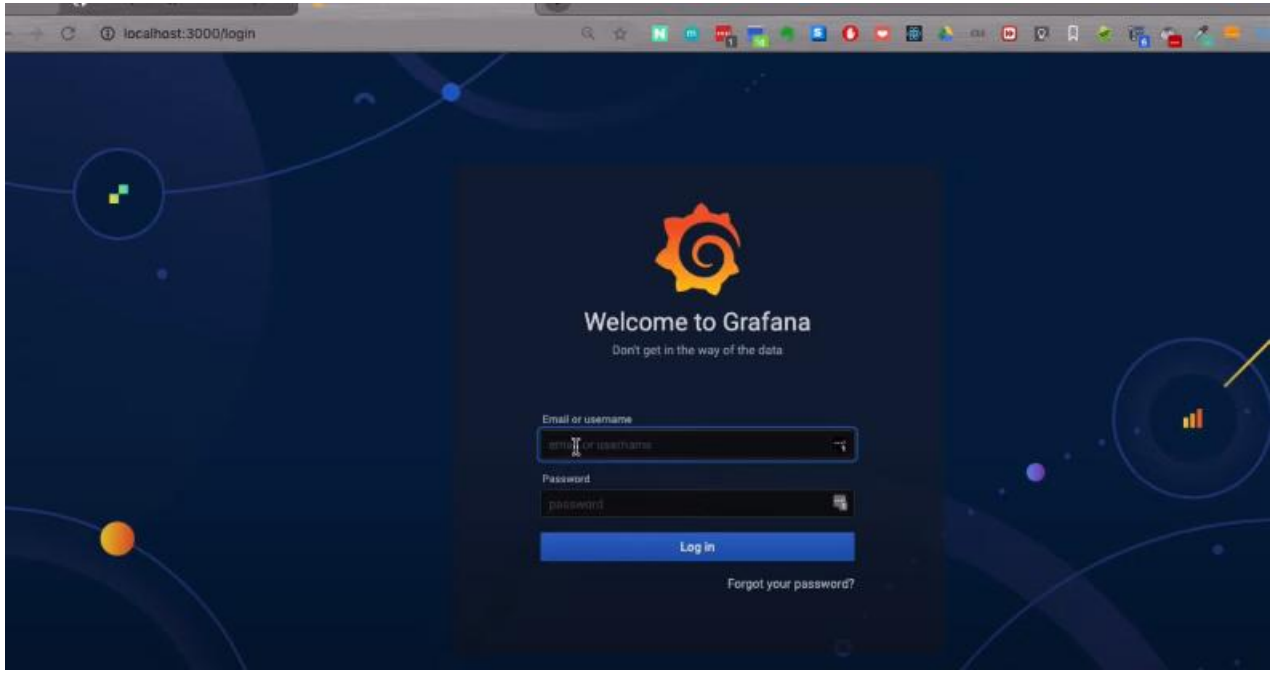

Grafana UI

Alert rules: shows the different alert rules

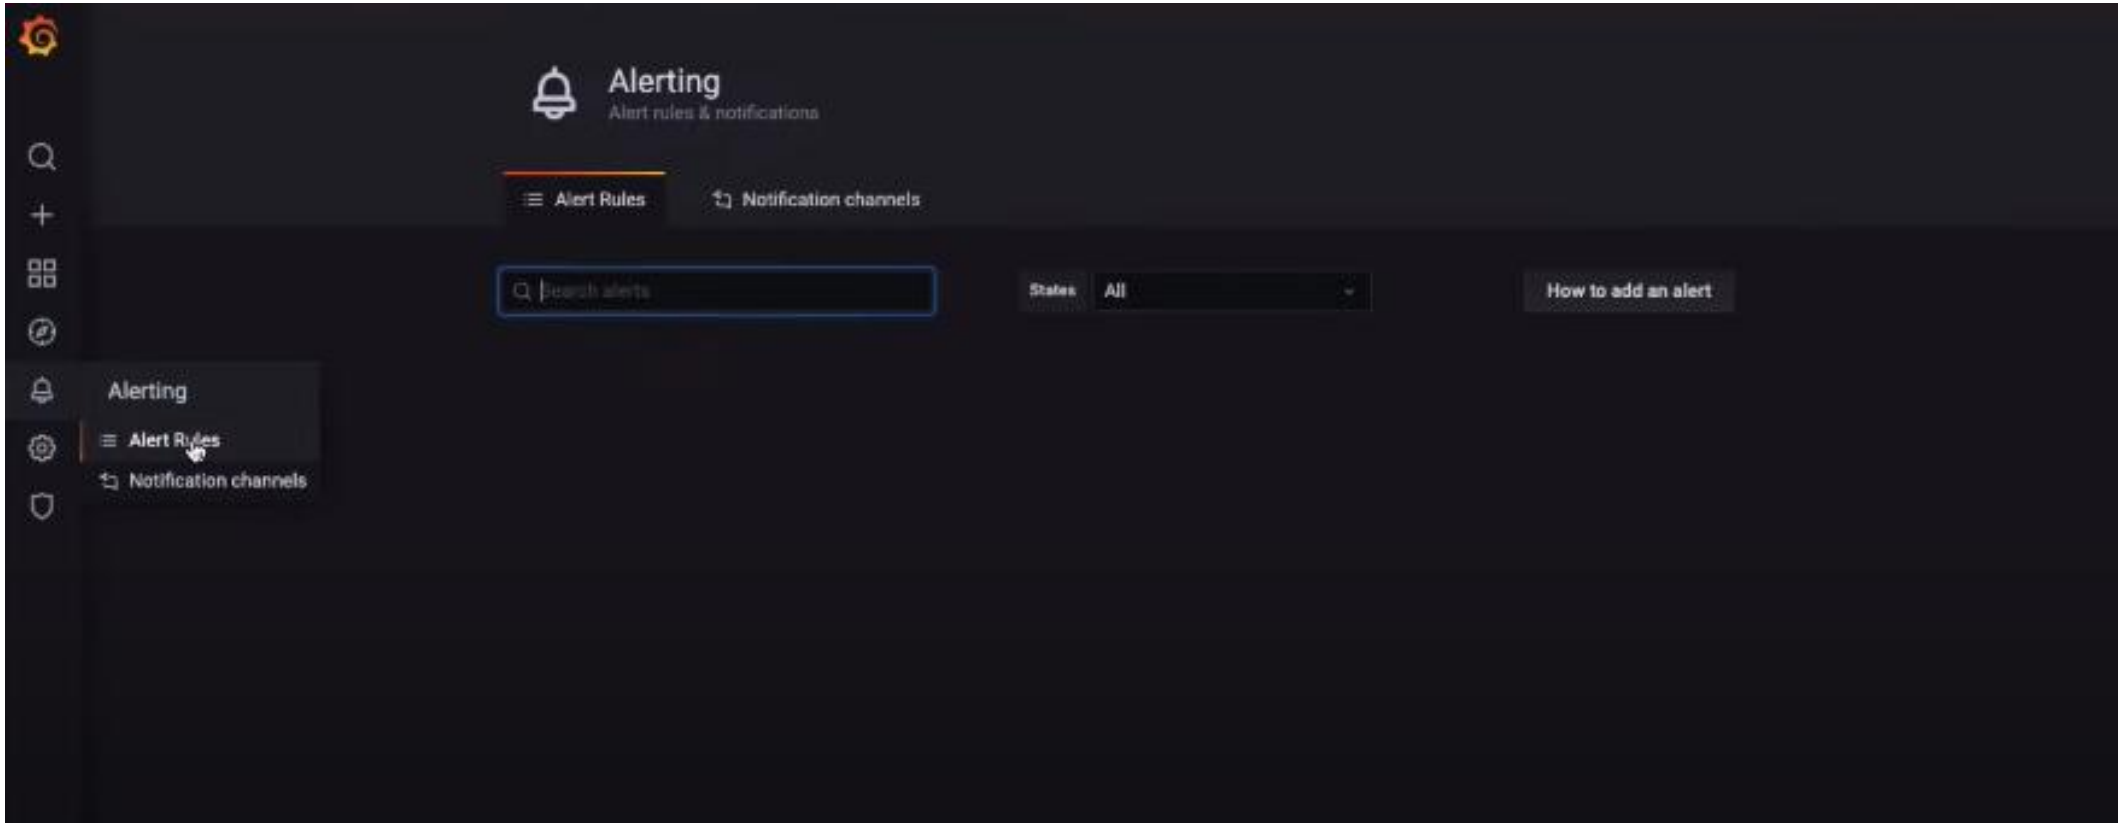

Grafana UI

PromQL: is the query language to create the data from Prometheus

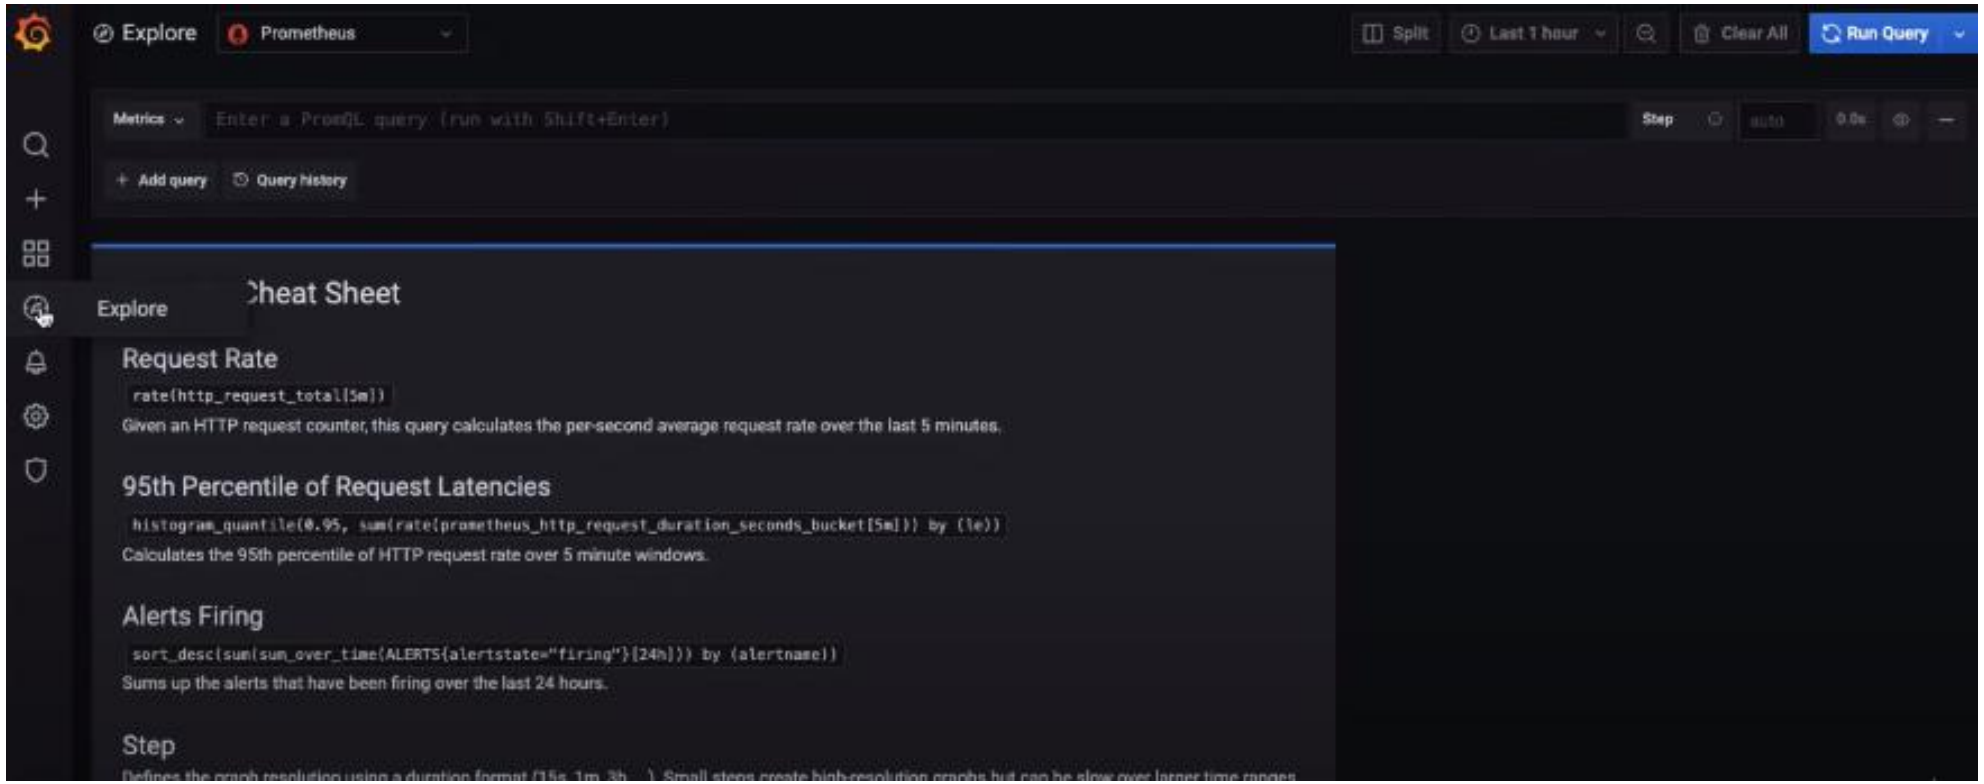

Grafana UI

Dashboard: shows the data that Prometheus is already collecting.

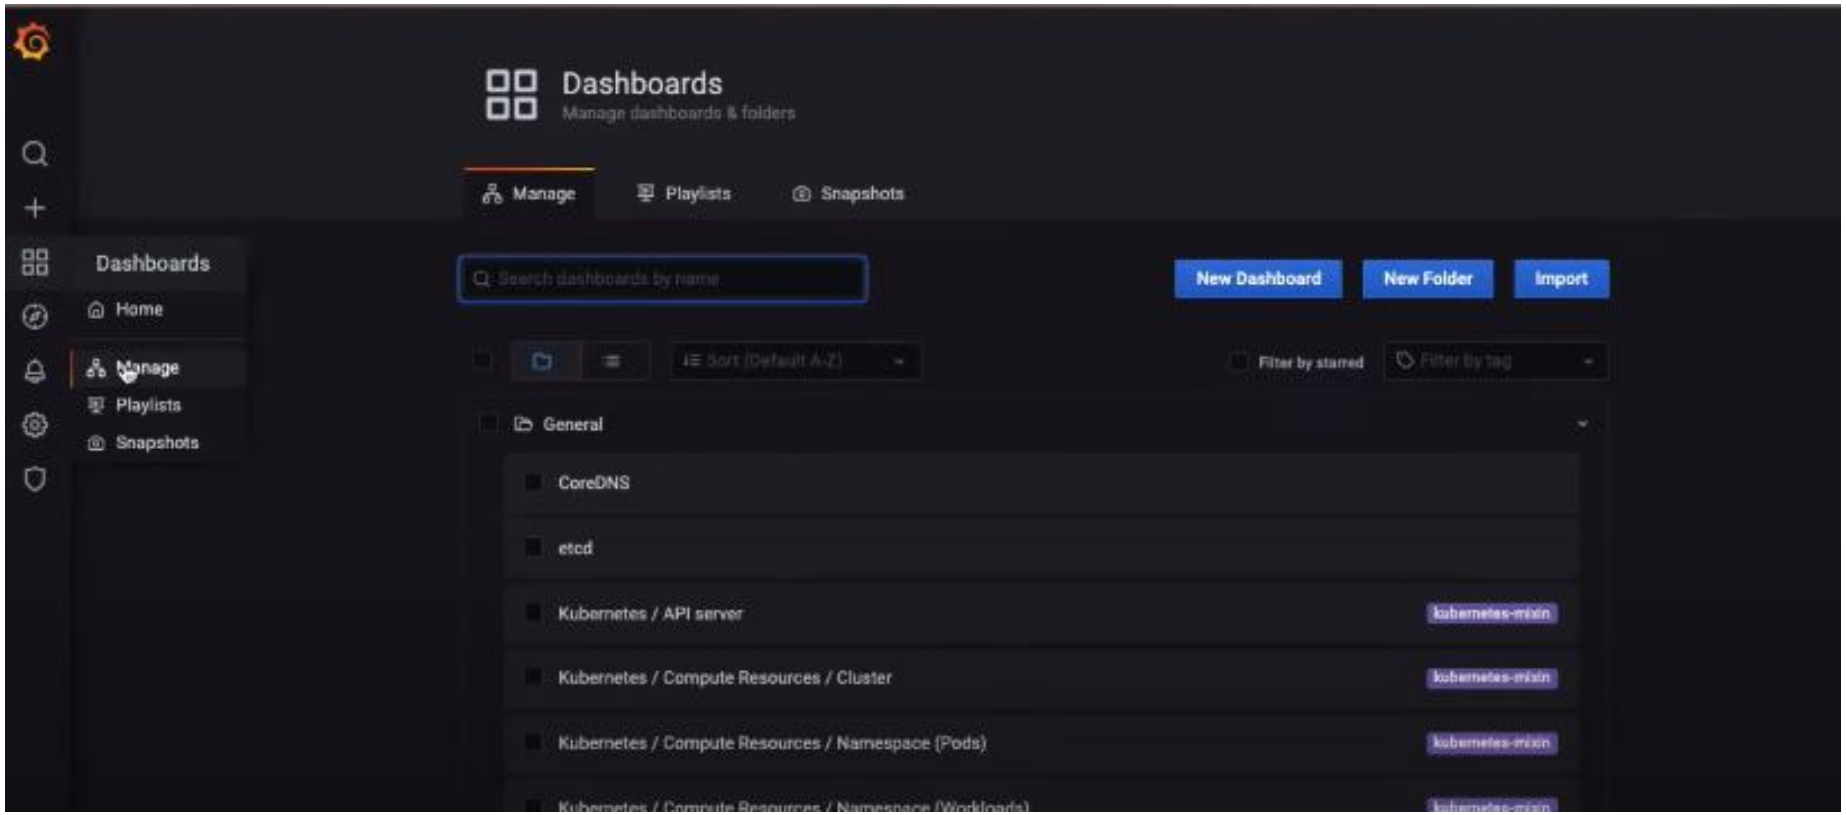

Grafana UI

Node exporter: shows the information about the nodes.

Grafana UI

Pod: shows the information about the pods.

Access Prometheus UI

Access the Prometheus using port forward.

```
[~]$ kubectl port-forward prometheus-prometheus-prometheus-oper-prometheus-0 9090
Forwarding from 127.0.0.1:9090 -> 9090
Forwarding from [::1]:9090 -> 9090
```

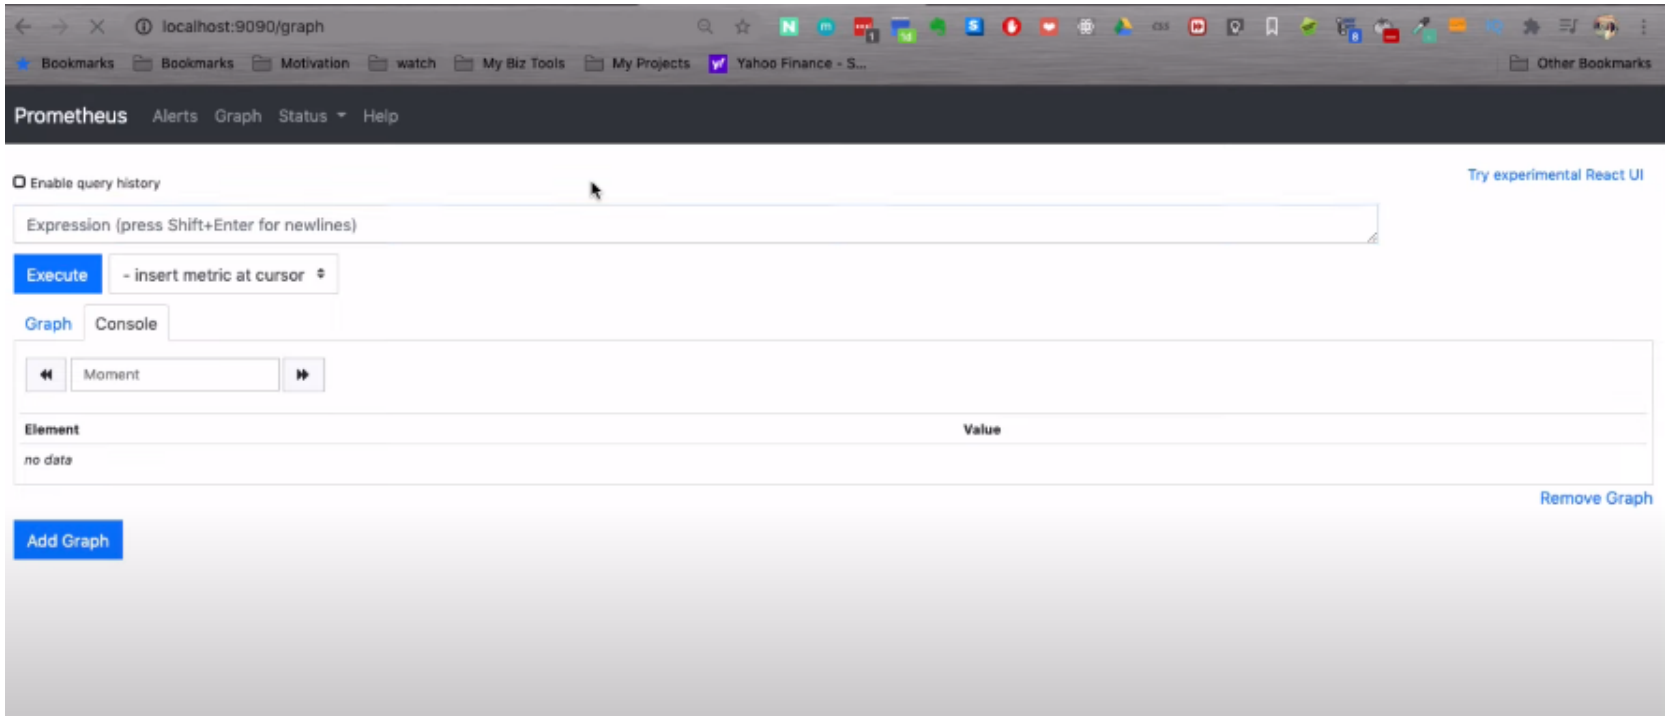

References

- https://prometheus.io/docs/introduction/overview/
- https://grafana.com/
- https://helm.sh/docs/

universität innsbruck

Thank you for your attention☺

Dr. Zahra Najafabadi Samani

Email: Zahra.Najafabadi-Samani@uibk.ac.at